CDS6214

Data Science Fundamentals

Lecture 2

Data Science Pipeline | Sources | Storage

data science pipelines

Data Science "Pipeline" Models

American data visualization expert

Ben Fry's Model (Data Visualization Process)

1. Acquire
2. Parse
3. Filter
4. Mine
5. Represent
6. Refine
7. Interact

Source: https://bcourses.berkeley.edu/courses/1377158/pages/cs-194-16-introduction-to-data-science-fall-2015

Data Science "Pipeline" Models

Data scientist and Co-founder of Cloudera

Jeff Hammerbacher's Model

1. Identify problem
2. Instrument data sources
3. Collect data
4. Prepare data (integrate, transform, clean, filter, aggregate)
5. Build model
6. Evaluate model
7. Communicate results

Source: https://bcourses.berkeley.edu/courses/1377158/pages/cs-194-16-introduction-to-data-science-fall-2015

Data Science Process

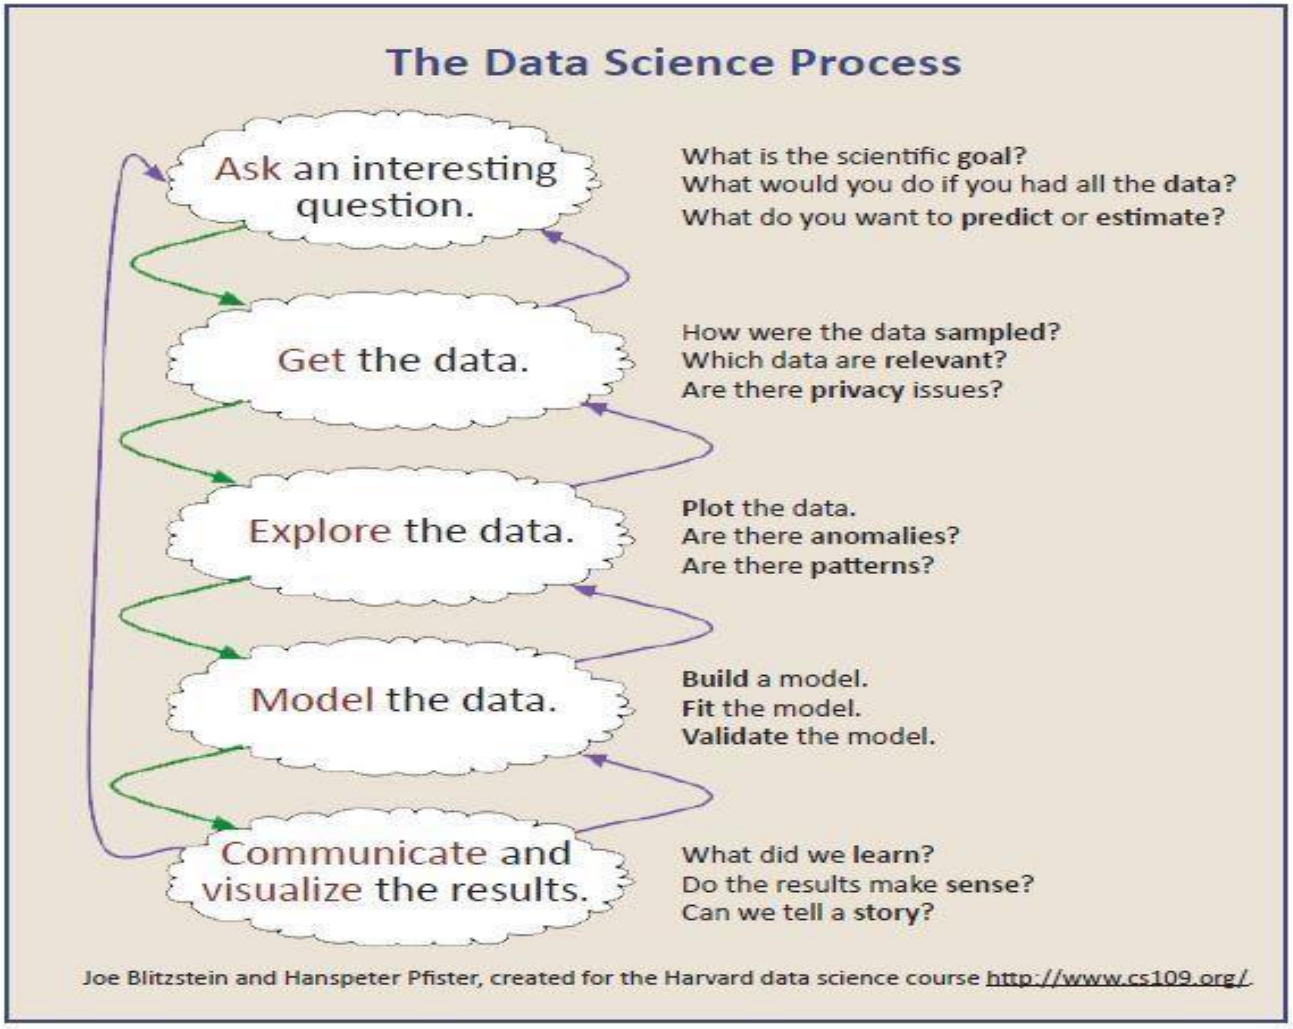

Process in-a-nutshell

Digging Around in Data

Clean, prep

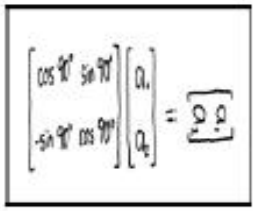

Hypothesize Model

Large Scale Exploitation

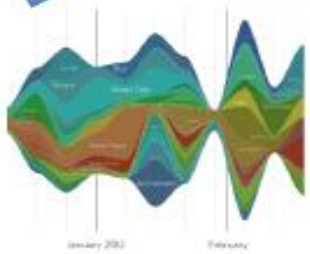

Evaluate Interpret

Source: https://bcourses.berkeley.edu/courses/1377158/pages/cs-194-16-introduction-to-data-science-fall-2015

data science pipelines

Data Science Process

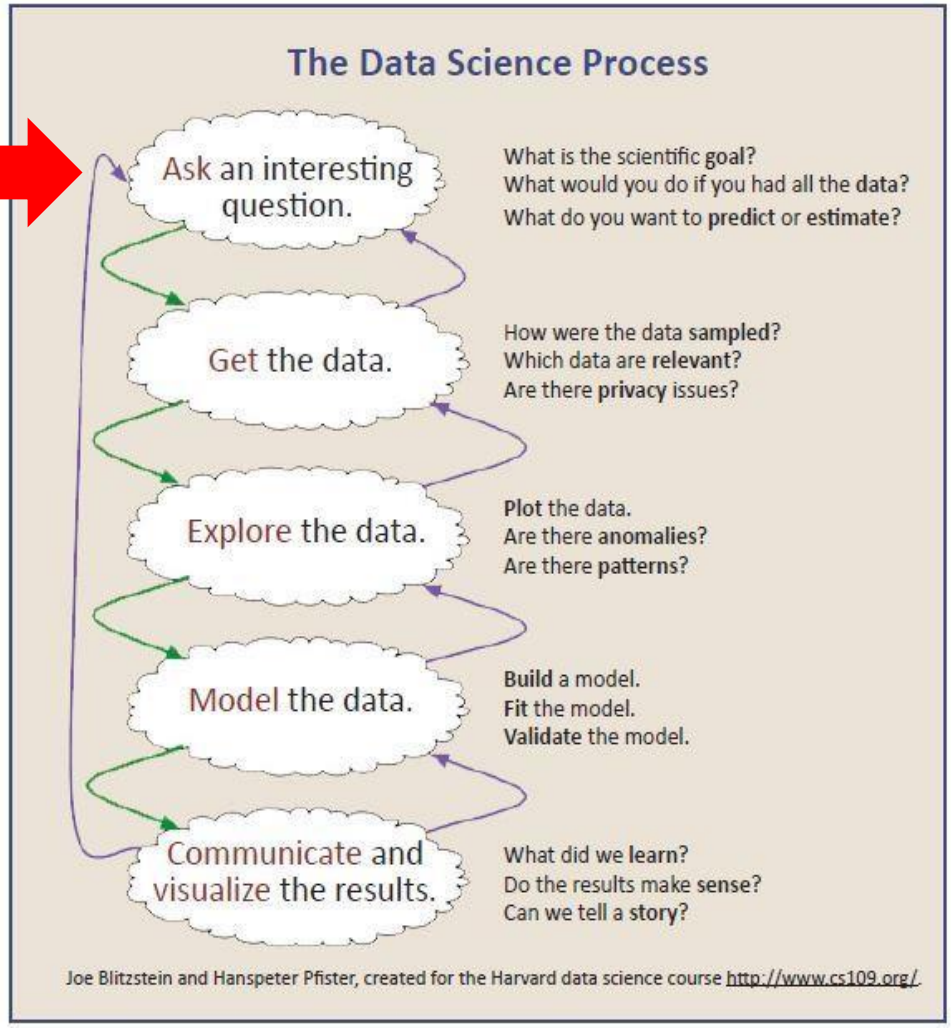

What should we do next?

- Identify the question
- Collect and pre-process the data
- Explore and analyze the data
- Model the data
- Infer and visualize results

The Question

There are 6 basic types of questions that can be asked:

1. **Descriptive**: A ***descriptive*** question is one that seeks to summarize a characteristic of a set of data.

   e.g. What is the mean number of servings of fresh fruits and vegetables per day?

2. **Exploratory**: An ***exploratory*** question is one in which you analyze the data to see if there are patterns, trends, or relationships between variables.

   e.g. What is the relationship between a range of dietary factors and viral illnesses?

The Question (2)

3. **Inferential**: An ***inferential*** question would be a restatement of the proposed hypothesis as a question and would be answered by analyzing a different set of data. The proposed hypothesis is usually derived from an exploratory question

e.g. Given the proposed hypothesis, that the habit of eating at least 5 servings of fresh fruit and vegetables per day is linked with fewer viral illnesses per year based on a sample population of US adults. Is this hypothesis also true for Asian population?

4. **Predictive**: A ***predictive*** question would be one where you ask what are the set of predictors / factors for a particular behaviour.

e.g. What type of people will eat a diet high in fresh fruits and vegetables during the next year?

The Question (3)

5. **Causal**: A ***causal*** question asks about whether changing one factor will change another factor, on average, in a population.

> *e.g. Will an increase in consumption of fresh fruits and vegetables reduce the frequency of contracting viral illnesses?*

6. **Mechanistic**: A ***mechanistic*** question points to *how* a factor affects the outcome.

> *e.g. How a diet high in fresh fruits and vegetables leads to a reduction in the number of viral illnesses.*

Characteristics of a Good Question

There are 5 basic characteristics of a good question:

1. The question should be of **interest** to your audience.
2. The question has **not already been answered.**
3. The question should also stem from a **plausible (valid correlations)** framework.
4. The question, should also, of course, be **answerable**.
5. The question should be **specific.**

Example

Is there a relation between the number of chicken burgers consumed per month and obesity among teens ?

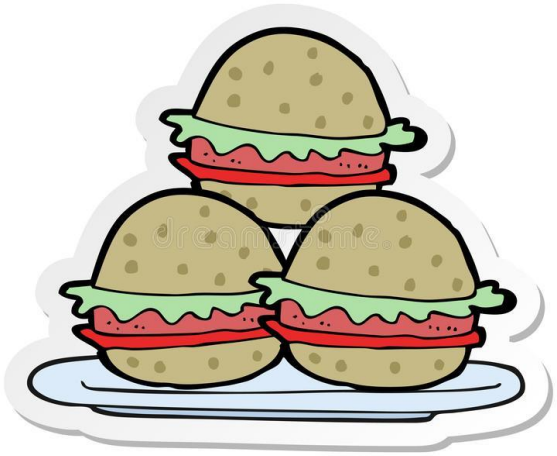

Characteristics of a Good Question

There are 5 basic characteristics of a good question:

1. The question should be of **interest** to your audience.
2. The question has **not already been answered.**
3. The question should also stem from a **plausible (valid correlations)** framework.
4. The question, should also, of course, be **answerable**.
5. The question should be **specific.**

Translating a Question into a Data Problem

- Every question must be operationalized as a **data problem** that leads to a result.
- Specifically, the questions asked should lead to **interpretable results.**
- Ensure that the data available to answer your question provide reasonably specific **measures** of the factors required to answer your question.
- **Potential problems:**
  - Confounding factor is a potential problem when your question asks about the relationship between factors.
  - Inappropriate data is used. The result is not interpretable because the underlying way in which the data was collected led to a biased result.

data science pipelines

Data Science Process

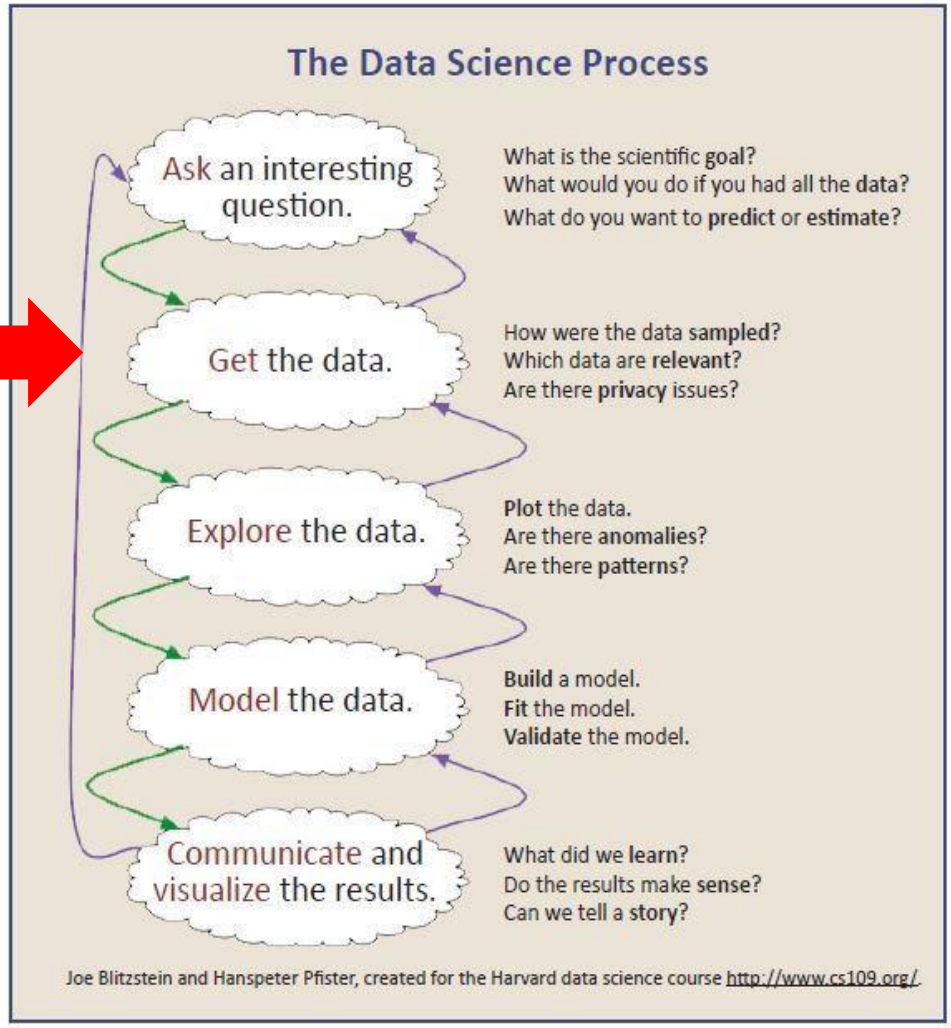

What should we do next?

- Identify the question
- Collect and pre-process the data
- Explore and analyze the data
- Model the data
- Infer and visualize results

What are the possible data sources?

- Generate the Data (surveys, interviews, etc.)
- Download the Data as files
- Connect to existing databases
- Call a web service API
- Setup equipment such as AV units and sensors
- Search for the Data (open data, search engines, industries/institutions)

What are the possible data sources?

- Generate the Data (surveys, interviews, etc.)
- Download the Data as files
- Connect to existing databases
- Call a web service API
- Setup equipment such as AV units and sensors
- Search for the Data (open data, search engines, industries/institutions)

Big Data Sources

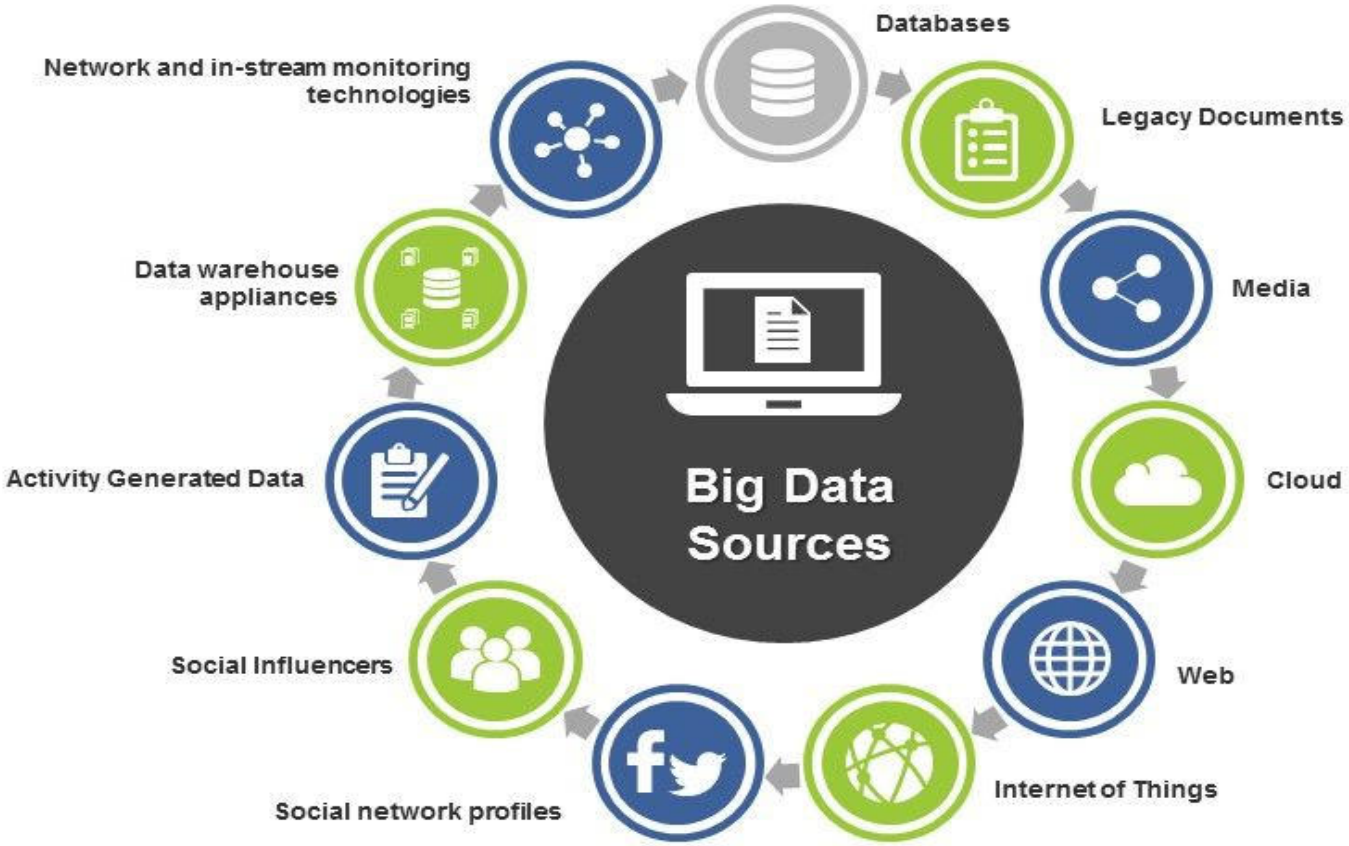

This slide is 100% editable. Adapt it to your needs and capture your audience's attention.

Data Collection

- Process of gathering and measuring information on targeted variables in an established systematic fashion, which then enables one to answer relevant questions and evaluate outcomes.
- Data can be *qualitative* or *quantitative*. Any difference between the two in terms of:
  a) Content?
  b) Methods of Collection?
  c) Example?

Common Problems in Data Collection

1. Irrelevant or duplicate data collected
2. Pertinent data omitted
3. Erroneous or misinterpreted data collected
4. Too little data acquired from client
5. Poor documentation
6. Conflicting data
7. Handwriting
8. Language barrier
9. Insufficient time

Common Data Files

Text, binary etc. – Some require more processing than others

- Syslog files
- Spreadsheets (xls, csv)
- PDF files
- Image files
- Raw text files
- Formatted text files

Common Data Format

- **Delimited values: Comma separated values, Tab separated values**
- Markup languages: HTML5 / XML, JSON
- Ad-hoc formats: Graph edge lists, voting records, device-captured signals, server logs, etc.

What is the **relation** between rows and columns?

```
Title,Author,ISBN13,Pages
1984,George Orwell,978-0451524935,268
Animal Farm,George Orwell,978-0451526342,144
Brave New World,Aldous Huxley,978-0060929879,288
Fahrenheit 451,Ray Bradbury,978-0345342966,208
Jane Eyre,Charlotte Brontë,978-0142437209,532
Wuthering Heights,Emily Brontë,978-0141439556,416
Agnes Grey,Anne Brontë,978-1593083236,256
Walden,Henry David Thoreau,978-1420922615,156
Walden Two,B. F. Skinner,978-0872207783,301
"Eats, Shoots & Leaves",Lynne Truss,978-1592400874,209
```

Tabular Data

What is a table?

- A **table** is a collection of **rows** and **columns**
- Each row has an **index**
- Each column has a **name**
- A **cell** is specified by an (index, name) pair
- A cell may or may not have a **value**

Often stored as text files in **CSV** or **TSV** format.

Tabular Data – Example

Fortune 500 Companies 2017 Data (open in spreadsheet editor)

	A	B	C	D	E	F	G	H	I	J	K
1		rank	name	employees	previousrank	revenues	revenuechange	profits	profitschange	assets	marketvalue
2	2	1	Walmart	2,300,000	1	$485,873	0.80%	$13,643.0	-7.20%	$198,825	$218,619
3	3	2	Berkshire Hathaway	367,700	4	$223,604	6.10%	$24,074.0	0.00%	$620,854	$411,035
4	4	3	Apple	116,000	3	$215,639	-7.70%	$45,687.0	-14.40%	$321,686	$753,718
5	5	4	Exxon Mobil	72,700	2	$205,004	-16.70%	$7,840.0	-51.50%	$330,314	$340,056
6	6	5	McKesson	68,000	5	$192,487	6.20%	$2,258.0	53.00%	$56,563	$31,439
7	7	6	UnitedHealth Group	230,000	6	$184,840	17.70%	$7,017.0	20.70%	$122,810	$157,793
8	8	7	CVS Health	204,000	7	$177,526	15.80%	$5,317.0	1.50%	$94,462	$81,310
9	9	8	General Motors	225,000	8	$166,380	9.20%	$9,427.0	-2.70%	$221,690	$52,968
10	10	9	AT&T	268,540	10	$163,786	11.60%	$12,976.0	-2.80%	$403,821	$255,679
11	11	10	Ford Motor	201,000	9	$151,800	1.50%	$4,596.0	-37.70%	$237,951	$46,349
12	12	11	AmerisourceBergen	18,500	12	$146,850	8.00%	$1,427.9	-	$33,656	$19,229
13	13	12	Amazon.com	341,400	18	$135,987	27.10%	$2,371.0	297.80%	$83,402	$423,031
14	14	13	General Electric	295,000	11	$126,661	-9.80%	$8,831.0	-	$365,183	$259,520
15	15	14	Verizon	160,900	13	$125,980	-4.30%	$13,127.0	-26.60%	$244,180	$198,900
16	16	15	Cardinal Health	37,300	21	$121,546	18.50%	$1,427.0	17.40%	$34,122	$25,725
17	17	16	Costco	172,000	15	$118,719	2.20%	$2,350.0	-1.10%	$33,163	$73,606
18	18	17	Walgreens Boots Alliance	300,000	19	$117,351	13.40%	$4,173.0	-1.10%	$72,688	$89,645
19	19	18	Kroger	443,000	17	$115,337	5.00%	$1,975.0	-3.10%	$36,505	$26,961

Tabular Data (csv)

Fortune 500 Companies 2017 Data (CSV format)

```
rank,company,revenue ($ millions),profit ($ millions)
1,Walmart,485873,13643
2,Berkshire Hathaway,223604,24074
3,Apple,215639,45687
4,Exxon Mobil,205004,7840
5,McKesson,192487,2258
6,UnitedHealth Group,184840,7017
7,CVS Health,177526,5317
8,General Motors,166380,9427
9,AT&T,163786,12976
10,Ford Motor,151800,4596
11,AmerisourceBergen,146850,1427.9
12,Amazon.com,135987,2371
13,General Electric,126661,8831
14,Verizon Communications,125980,13127
15,Cardinal Health,121546,1427
16,Costco,118719,2350
17,Walgreens Boots Alliance,117351,4173
18,Kroger,115337,1975
19,Chevron,107567,-497
20,Fannie Mae,107162,12313
21,J.P. Morgan Chase,105486,24733
22,Express Scripts Holding,100288,3404.4
23,Home Depot,94595,7957
24,Boeing,94571,4895
25,Wells Fargo,94176,21938
26,Bank of America Corp.,93662,17906
27,Alphabet,90272,19478
28,Microsoft,85320,16798
```

Common Data Formats

- Delimited values: Comma separated values, Tab separated values
- **Markup languages: HTML5 / XML, JSON**
- Ad-hoc formats: Graph edge lists, voting records, device-captured signals, server logs, etc.

Example of markup language:HTML5

```
<!DOCTYPE html><html>
<head>
<title>HTML 5 Demo</title><style>
.DSF {
font-size:40px;font-weight:bold;color:green;
}body {
text-align:center;}
</style></head><body>
<div class = "DSF">DataScienceFundamental</div><aside>
<div>A computer science portal for geeks</div>
</aside></body></html>
```

Source: https://www.geeksforgeeks.org/html5/

Example of markup language: XML

XML: Generalizes HTML and specifies data **structure**. XML schema can be applied later to interpret XML data and specify **data types**. Here is a sample XML-encoded **data**:

[Example]

```
<location>
  <latitude>37.78333</latitude>
  <longitude>122.4167</longitude>
</location>
```

When stored without a schema, the numerical data are stored as **strings**.

Source: https://bcourses.berkeley.edu/courses/1267848

JSON (JavaScript Object Notation)

- A lightweight data-interchange format.
- It is based on a subset of the JavaScript Programming Language, Standard ECMA-262 3rd Edition - December 1999.
- Completely language independent but uses conventions that are familiar to programmers of the C-family of languages, including C, C++, C#, Java, JavaScript, Perl, Python, and many others.
- Consists of two components:
  - **A collection of name/value pairs**. In various languages, this is realized as an object, record, struct, dictionary, hash table, keyed list, or associative array.
  - **An ordered list of values.** In most languages, this is realized as an array, vector, list, or sequence.

JSON – Example Tweet Format


```
{"location":[
        {"latitude":37.78333, "longitude":122.4167 }
]}
```

Common Data Formats

- Delimited values: Comma separated values, Tab separated values
- Markup languages: HTML5 / XML, JSON,
- **Ad-hoc formats: Graph edge lists, voting records, device-captured signals, server logs, etc.**

Example – Graph Edge Lists

Let's look at definition of a graph extracted exactly as it is from a website:

A graph is a data structure that consists of a finite set of vertices, which are also called nodes, and a set of edges, which are references/links between the vertices. The edges of a graph are represented as ordered or unordered pairs, depending on whether or not the graph is directed or undirected.

https://medium.com/basecs/from-theory-to-practice-representing-graphs-cfd782c5be38

Example – Graph Edge Lists

Figure A

Figure B

For each figure, can you identify the set of vertices and edges ?

Example – Graph Edge Lists

An edge list is a list (or array) of all of the |E| edges in a graph. Edge lists are one of the simplest representations of

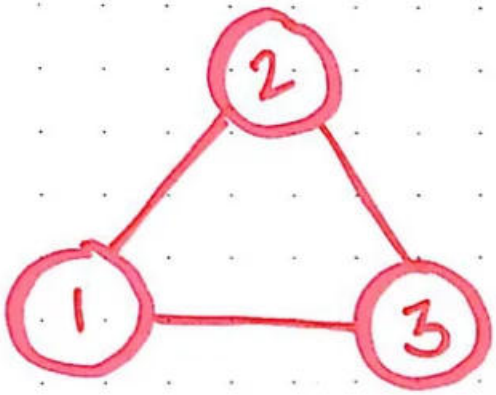

To represent a graph with three nodes (1, 2, and 3), we could either use a list format or an array to represent the graph as an edge list.

https://medium.com/basecs/from-theory-to-practice-representing-graphs-cfd782c5be38

Example – Graph Edge Lists

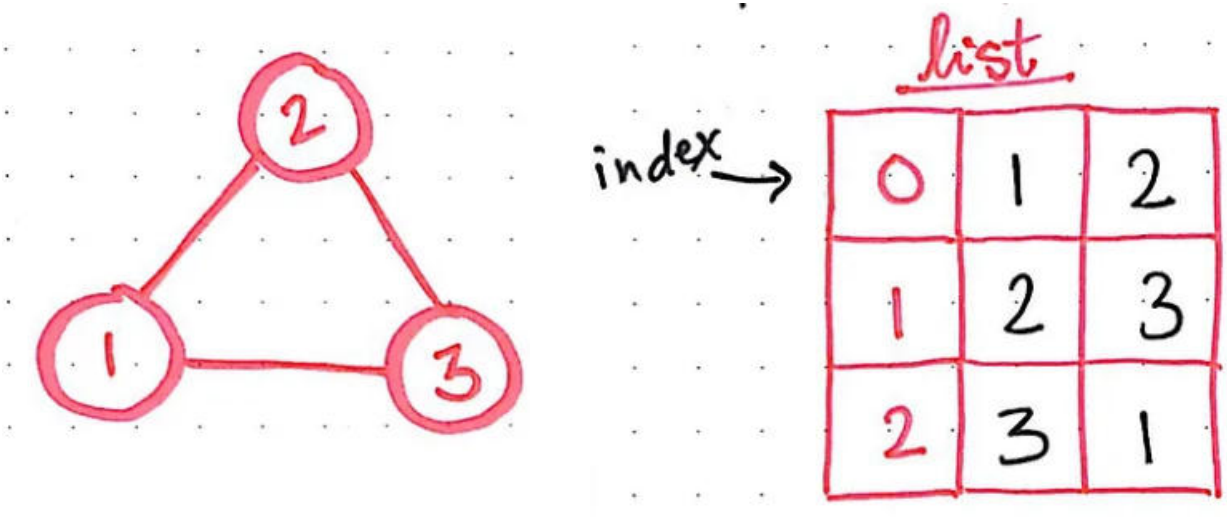

How do you represent the above graph as an array ?

Example – Apache Web Log

Processes, usually daemons, create logs
e.g., `httpd, mysqld, syslogd`

```
66.249.65.107 - - [08/Oct/2007:04:54:20 -0400] "GET /support.html
HTTP/1.1" 200 11179 "-" "Mozilla/5.0 (compatible; Googlebot/2.1;
+http://www.google.com/bot.html)"

111.111.111.111 - - [08/Oct/2007:11:17:55 -0400] "GET / HTTP/1.1"
200 10801 "http://www.google.com/search?q=log+analyzer&ie=utf-
8&oe=utf-8 &aq=t&rls=org.mozilla:en-US:official&client=firefox-a"
"Mozilla/5.0 (Windows; U; Windows NT 5.2; en-US; rv:1.8.1.7)
Gecko/20070914 Firefox/2.0.0.7"

111.111.111.111 - - [08/Oct/2007:11:17:55 -0400] "GET /style.css
HTTP/1.1" 200 3225 ""http://www.loganalyzer.net/" "Mozilla/5.0
(Windows; U; Windows NT 5.2; en-US; rv:1.8.1.7) Gecko/20070914
Firefox/2.0.0.7"
```

Example of setup for data collection for wearable and mobile technology

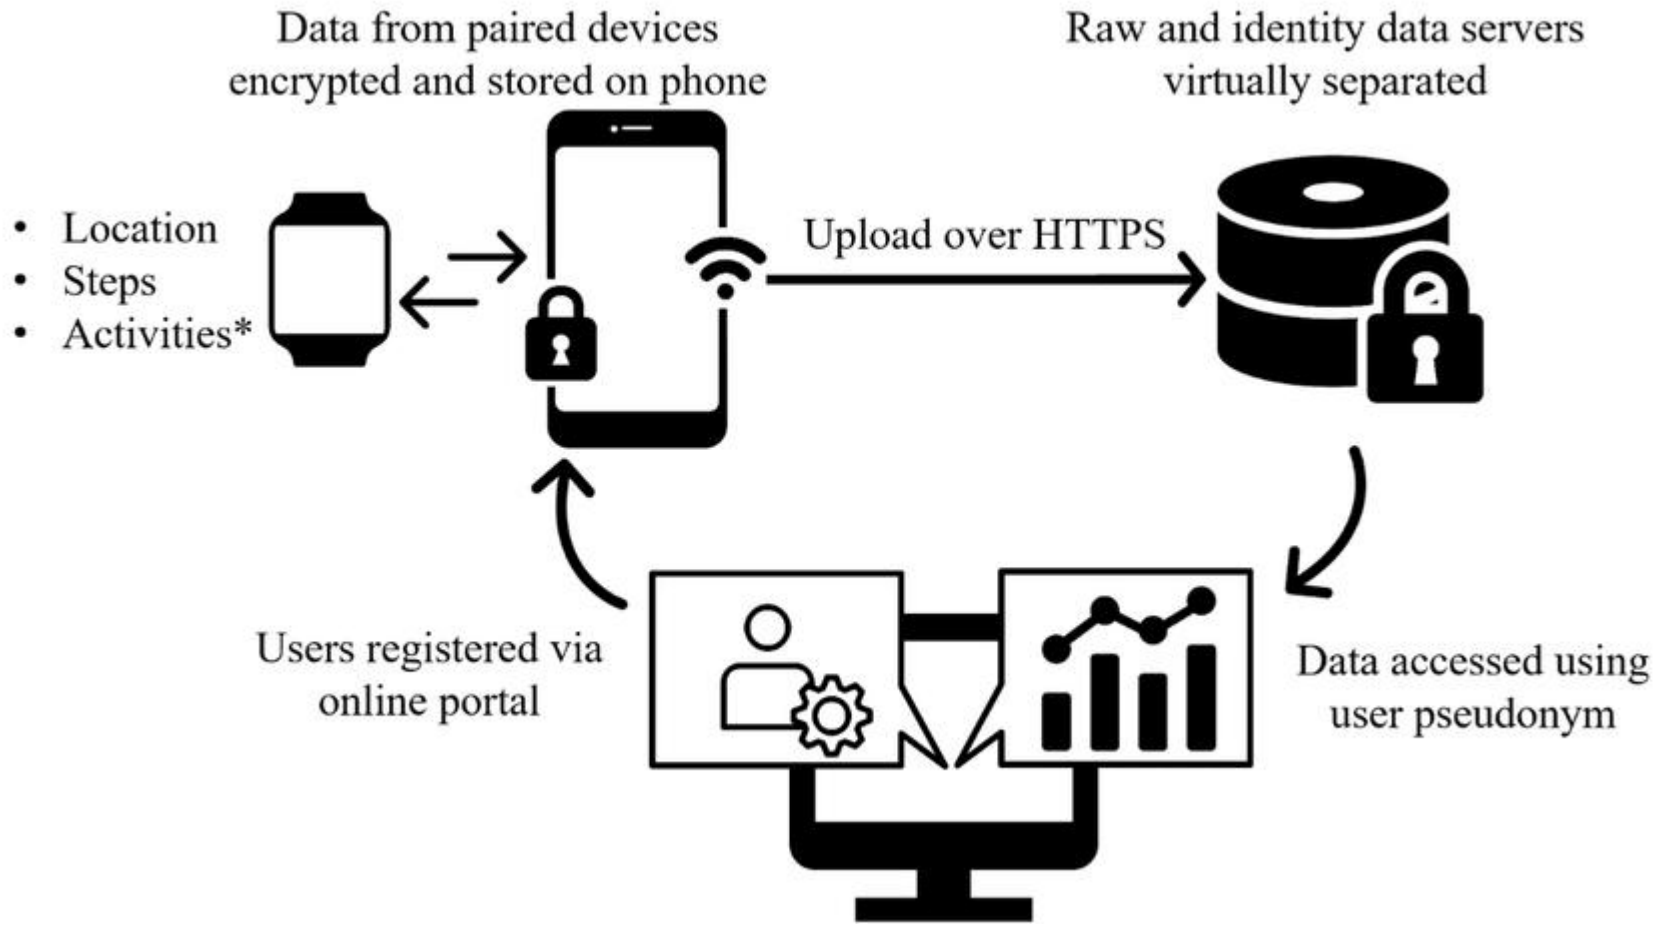

https://www.researchgate.net/figure/Data-collection-setup-data-is-collected-using-sensors-on-board-a-smartwatch-and_fig5_328379775

Time based Data – Device Measurements

2010-11-04	05:40:51.303739	M004	ON	Bed_to_Toilet	begin
2010-11-04	05:40:52.342105	M005	OFF		
2010-11-04	05:40:57.176409	M007	OFF		
2010-11-04	05:40:57.941486	M004	OFF		
2010-11-04	05:43:24.021475	M004	ON		
2010-11-04	05:43:26.273181	M004	OFF		
2010-11-04	05:43:26.345503	M007	ON		
2010-11-04	05:43:26.793102	M004	ON		
2010-11-04	05:43:27.195347	M007	OFF		
2010-11-04	05:43:27.787437	M007	ON		
2010-11-04	05:43:29.711796	M005	ON		
2010-11-04	05:43:30.279021	M004	OFF	Bed_to_Toilet	end
2010-11-04	05:43:45.7324	M003	ON	Sleeping	begin
2010-11-04	05:43:52.044085	M003	OFF		
2010-11-04	05:43:53.185335	M002	ON		
2010-11-04	05:43:53.253809	M003	ON		
2010-11-04	05:43:59.493281	M002	OFF		
2010-11-04	05:44:04.048766	M003	OFF		
2010-11-04	05:44:06.14204	M003	ON		
2010-11-04	05:44:11.229146	M003	OFF		

Sample raw and activity annotated sensor data.Sensors IDs starting with M are motion sensors

https://www.researchgate.net/figure/Sample-raw-and-activity-annotated-sensor-dataSensors-IDs-starting-with-M-are-motion_fig2_308843277

Malaysia's Official Open Data Portal

https://data.gov.my/

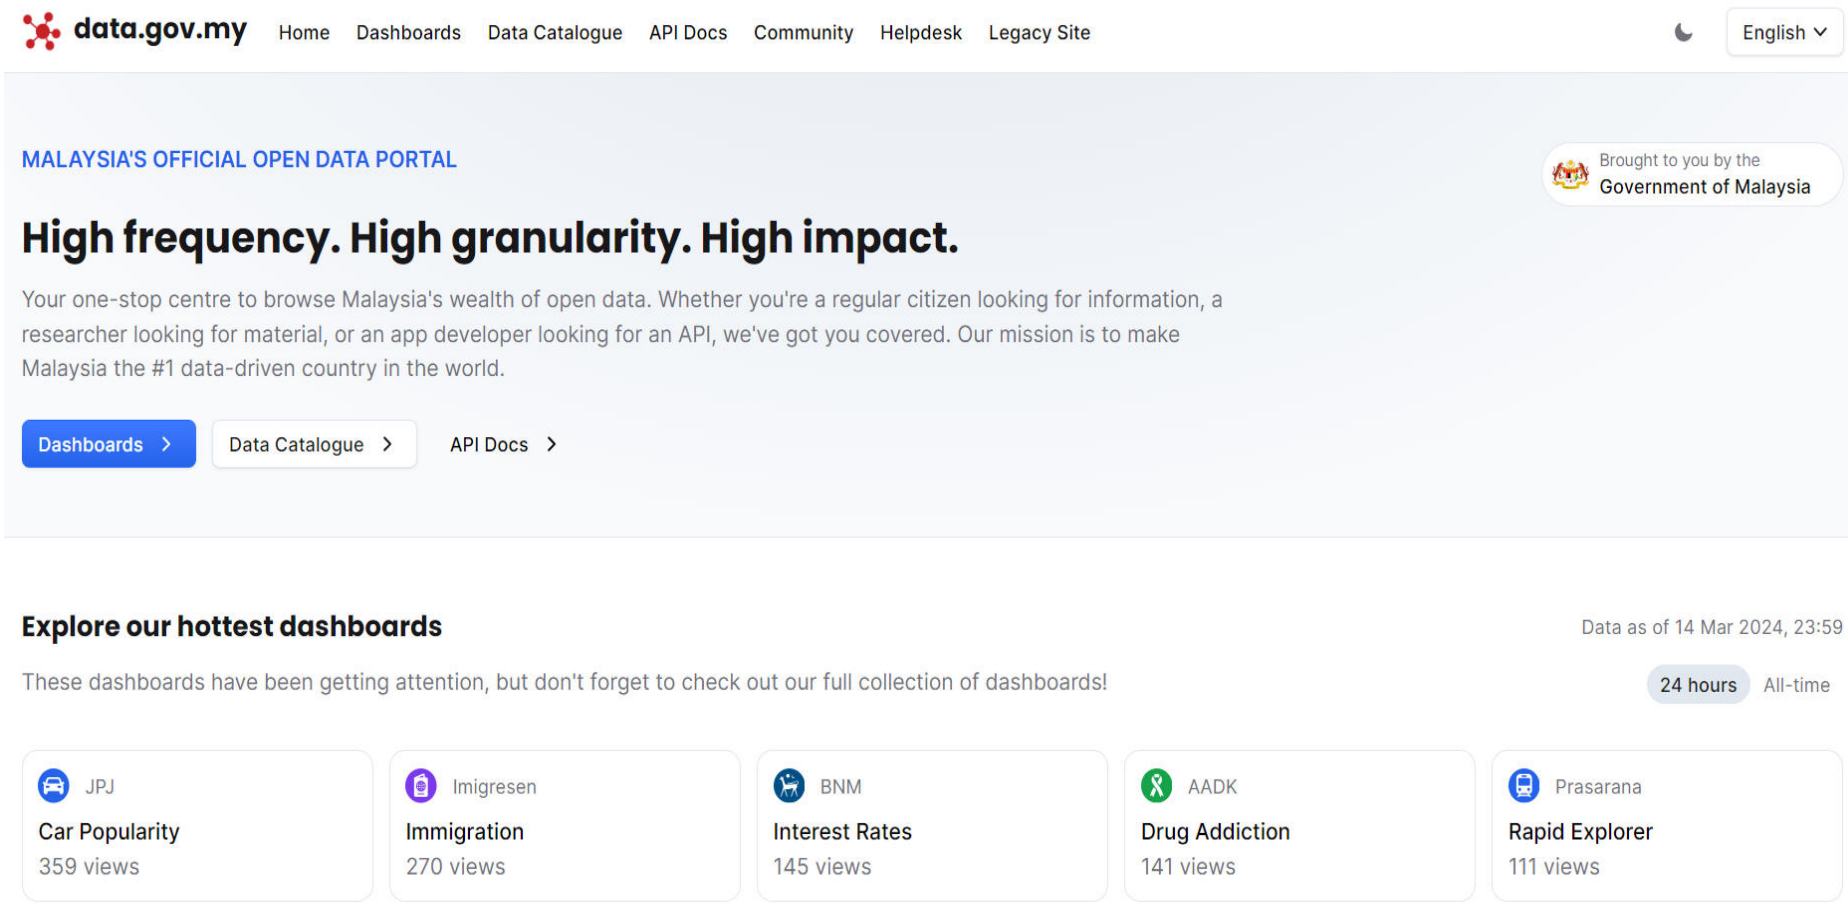

The Home of the U.S. Government's Open Data

https://data.gov/

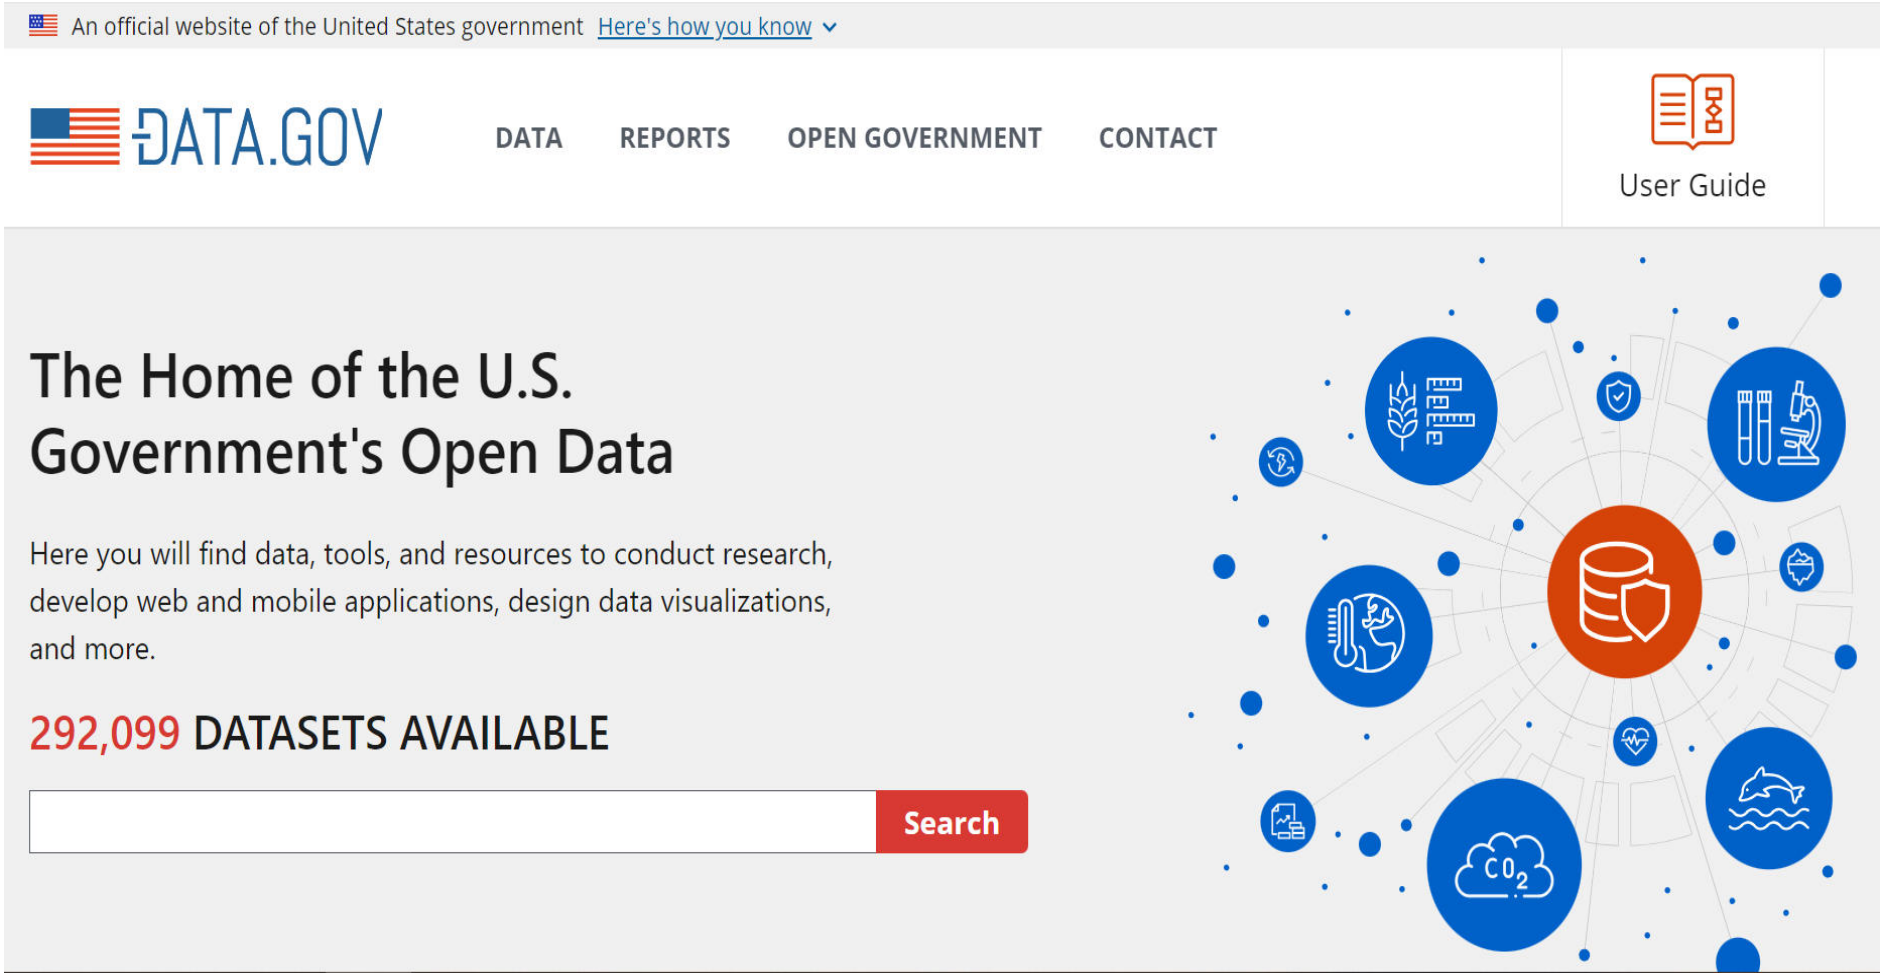

The Official Portal for European Data

https://data.europa.eu/en

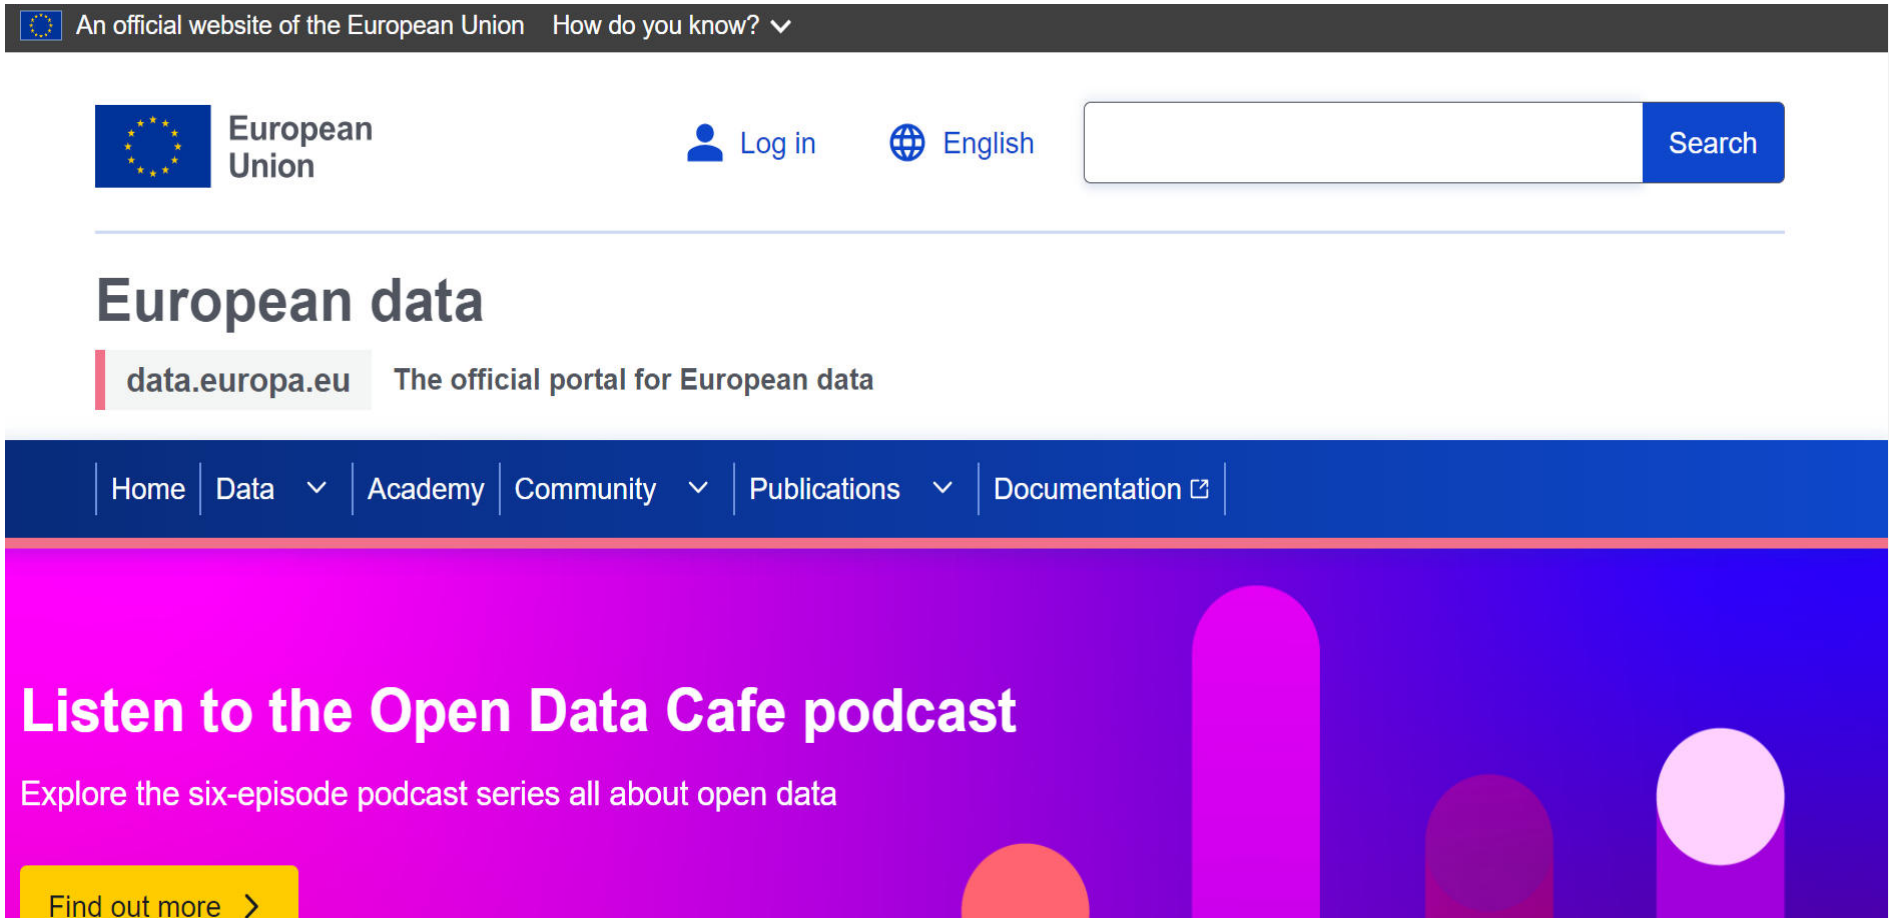

Google Dataset Search

https://datasetsearch.research.google.com/

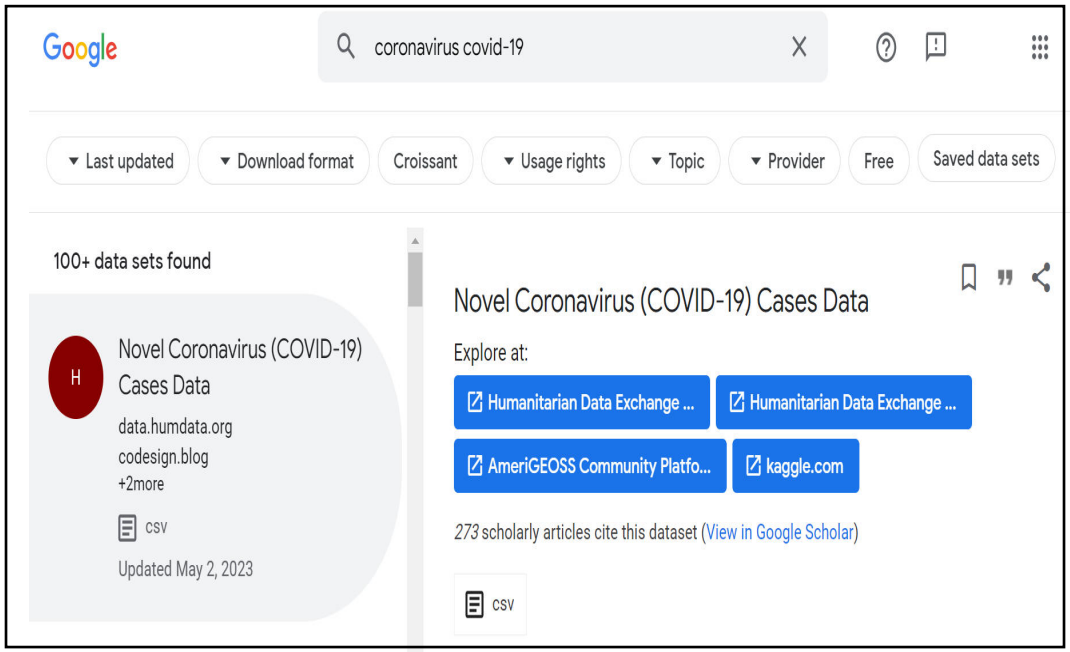

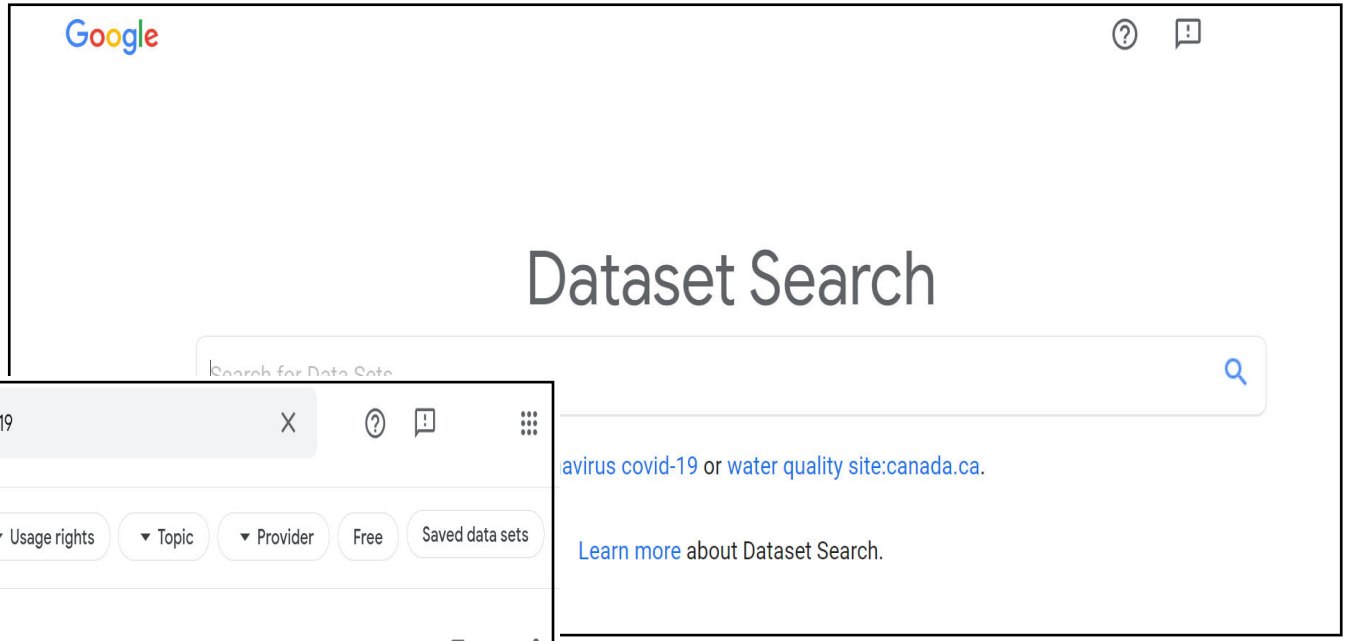

types of data store

a repository for persistently storing and managing collections of data which include not just repositories like databases, but also simpler store types such as simple files, emails etc.

By 2025,
an estimated 463 exabytes of data will be generated every single day
(1EB = 1 million TB)

Data Sources: Example of a Web Company

Example: Facebook

- Application databases
- Web server logs
- Event logs
- API server logs
- Ad server logs
- Search server logs
- Advertisement landing page content
- Wikipedia
- Images and video

Structured Data

Semi-structured Data

Unstructured Data

Source: https://bcourses.berkeley.edu/courses/1267848

The (changing) role of Schema

A data **schema** specifies the **structure** and **types** of a data repository, e.g. the types of each column in a table. They may also specify constraints **within** or **between** data fields.

Traditional databases (Relational DB) are **schema-on-write**. You cannot load data into a table without a schema.

Newer data stores (e.g. noSQL) are **schema-on-read** or **schema-less**: You can defer applying a schema until you read the data, or avoid schema altogether

Source: https://bcourses.berkeley.edu/courses/1267848

Relational DB

- **Relational databases** use the notion of databases separated into tables where each column represents a field and each row represents a record
  - Named after a branch of algebraic set theory known as relational algebra
- Tables can be related or linked with each other with use of foreign keys / common columns
- E.g. of Relational DBs: MySQL, PostgreSQL SQLite3
  - Represent and store data in tables and rows
- **Relational database management systems (RDBMSs)** are the primary technology for storing structured data in web and business applications

Relational DB Schema

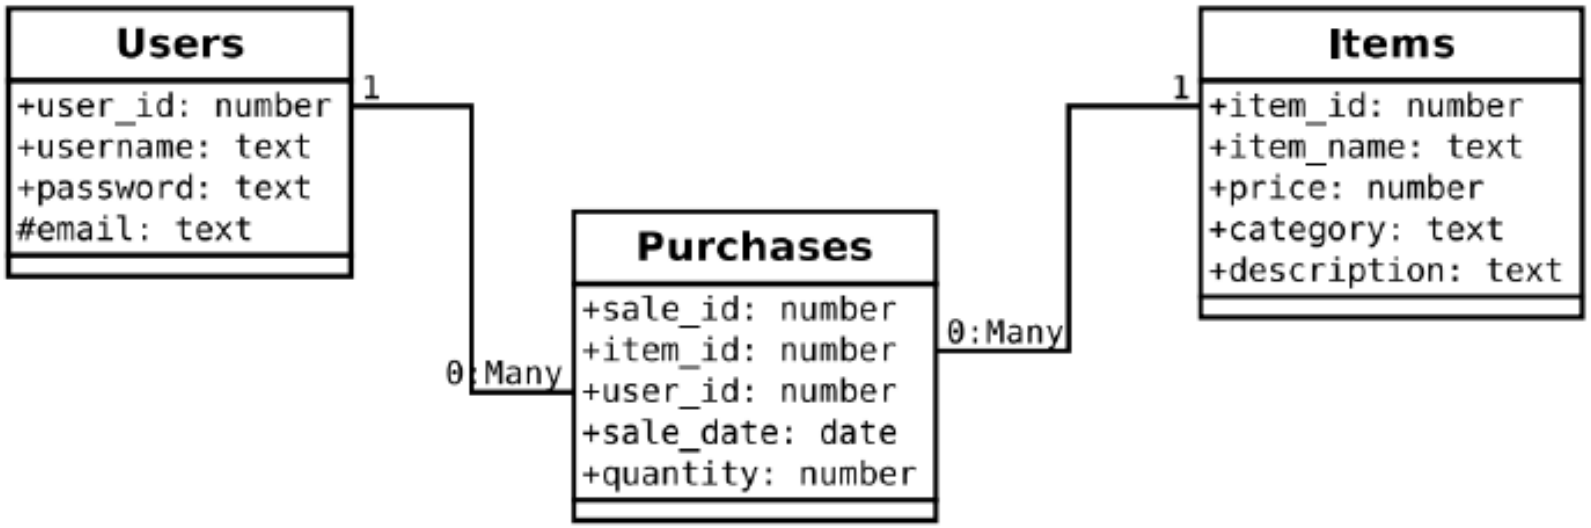

Relational DB

- An important design aspect: **Normalisation** of schema
  - Reduces redundancy (repetitive data), ensure data stored logically
  - Larger tables are divided into smaller tables which are linked using relationships

Normalisation of Schema

- 3 common forms:
  a. **First Normal Form (1NF)**: Eliminate groups of repeating data by creating a new table for each group of related data which is identified by primary key.
  b. **Second Normal Form (2NF)**: If a set of values are the same for multiple records, move them to a new table and link the two tables with foreign key.
  c. **Third Normal Form (3NF):** Fields which do not depend on the primary key of a table must be removed and if necessary be put into another table.

1NF

Sales Records:

Cust Name	Item	Shipping Address	Newsletter	Supplier	Supplier Phone	Price
Alan Smith	Xbox One	35 Palm St, Miami	Xbox News	Microsoft	(800) BUY-XBOX	250
Roger Banks	PlayStation 4	47 Campus Rd, Boston	PlayStation News	Sony	(800) BUY-SONY	300
Evan Wilson	Xbox One, PS Vita	28 Rock Av, Denver	Xbox News, PlayStation News	Wholesale	Toll Free	450
Alan Smith	PlayStation 4	47 Campus Rd, Boston	PlayStation News	Sony	(800) BUY-SONY	300

1st Normal Form
- Each cell to be Single valued
- Entries in a column are same type
- Rows uniquely identified - Add Unique ID, or Add more colums to make unique
  (Note: The order of the rows and the order of the columns are irrelevant)

Primary Key

Order_ID	Cust Name	Item	Shipping Address	Newsletter	Supplier	Supplier Phone	Price
1	Alan Smith	Xbox One	35 Palm St, Miami	Xbox News	Microsoft	(800) BUY-XBOX	250
2	Roger Banks	PlayStation 4	47 Campus Rd, Boston	PlayStation News	Sony	(800) BUY-SONY	300
3	Evan Wilson	Xbox One	28 Rock Av, Denver	Xbox News	Microsoft	(800) BUY-XBOX	250
4	Evan Wilson	PS Vita	28 Rock Av, Denver	PlayStation News	Sony	(800) BUY-SONY	200
5	Alan Smith	PlayStation 4	47 Campus Rd, Boston	PlayStation News	Sony	(800) BUY-SONY	300

Extracted and modified from source: https://ericdatascience.wordpress.com/2019/06/12/database-normalisation

2NF

Primary Key

Order_ID	Cust Name	Item	Shipping Address	Newsletter	Supplier	Supplier Phone	Price
1	Alan Smith	Xbox One	35 Palm St, Miami	Xbox News	Microsoft	(800) BUY-XBOX	250
2	Roger Banks	PlayStation 4	47 Campus Rd, Boston	PlayStation News	Sony	(800) BUY-SONY	300
3	Evan Wilson	Xbox One	28 Rock Av, Denver	Xbox News	Microsoft	(800) BUY-XBOX	250
4	Evan Wilson	PS Vita	28 Rock Av, Denver	PlayStation News	Sony	(800) BUY-SONY	200
5	Alan Smith	PlayStation 4	47 Campus Rd, Boston	PlayStation News	Sony	(800) BUY-SONY	300

2nd Normal Form - All attributes (Non-Key Columns) dependent on the key

Primary Key

Order_ID	Cust Name	Shipping Address	Newsletter
1	Alan Smith	35 Palm St, Miami	Xbox News
2	Roger Banks	47 Campus Rd, Boston	PlayStation News
3	Evan Wilson	28 Rock Av, Denver	Xbox News
4	Evan Wilson	28 Rock Av, Denver	PlayStation News
5	Alan Smith	47 Campus Rd, Boston	PlayStation News

Primary Key

Item	Supplier	Supplier Phone	Price
Xbox One	Microsoft	(800) BUY-XBOX	250
PlayStation 4	Sony	(800) BUY-SONY	300
PS Vita	Sony	(800) BUY-SONY	200

Primary Key | Foreign Key

Order_ID	Item
1	Xbox One
2	PlayStation 4
3	Xbox One
4	PS Vita
5	PlayStation 4

Extracted and modified from source: https://ericdatascience.wordpress.com/2019/06/12/database-normalisation

3NF

3rd Normal Form - All Fields (columns) can be determined Only by the Key in the table and and no other column

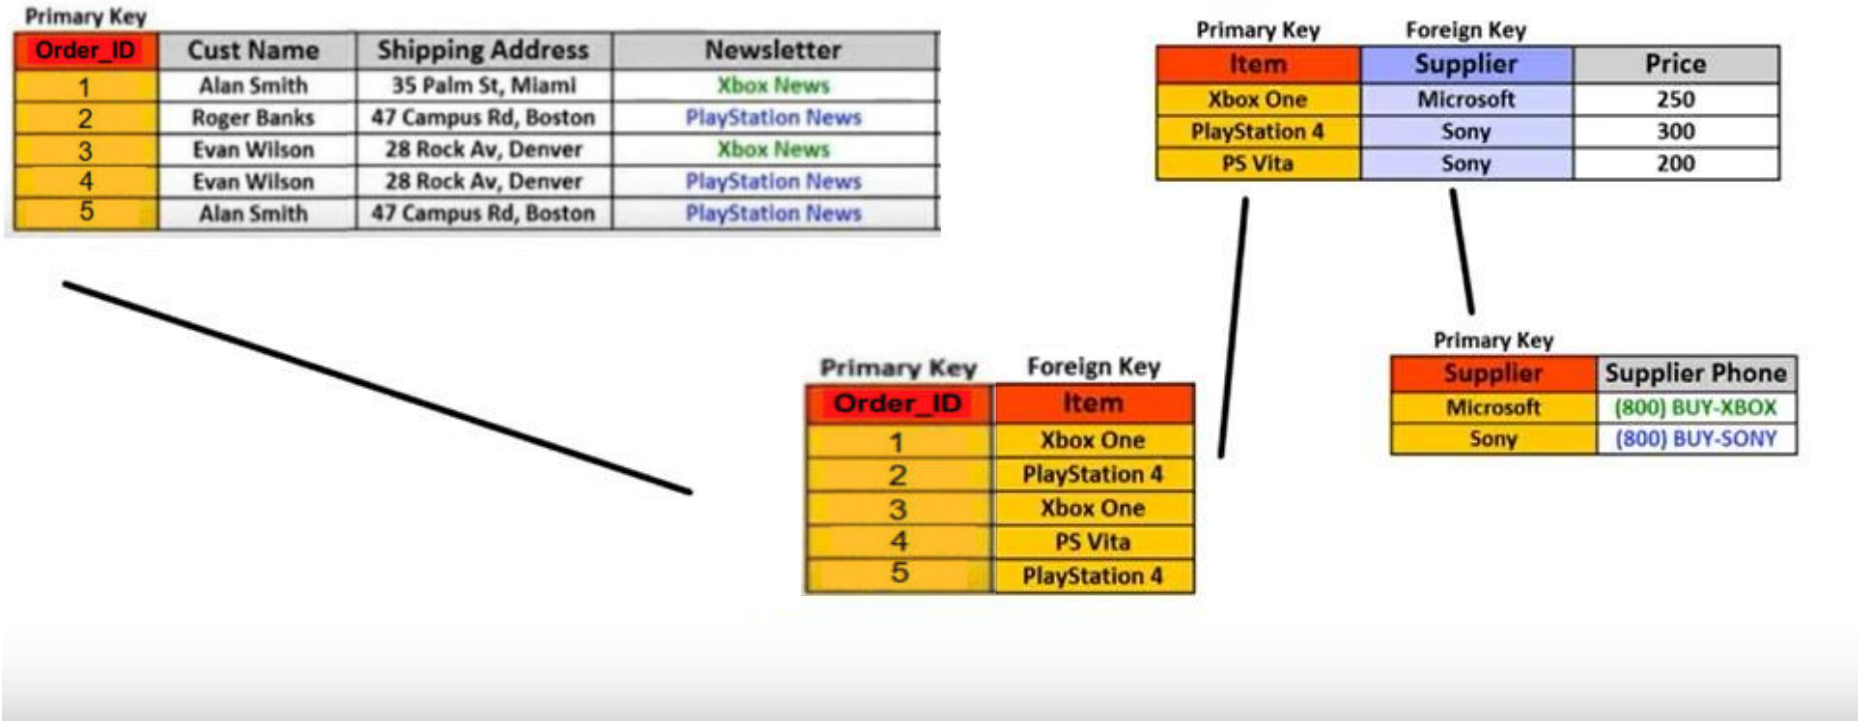

Extracted and modified from source: https://ericdatascience.wordpress.com/2019/06/12/database-normalisation

Relational DB

- An important aspect to guarantee reliability of transactions – adherence to the ACID properties:
  a. **Atomicity:** Either all parts of a transaction must be completed or none
  b. **Consistency:** The integrity of the database is preserved by all transactions. DB is not left in invalid state after a transaction
  c. **Isolation**: A transaction must be run isolated in order to guarantee inconsistency in data does not affect other transactions
  d. **Durability**: Changes made by a completed transaction must be preserved or durable

Limitations of Traditional RDBMS

- Traditional RDBMS are not feasible solutions to all data storage problems → obvious limits and difficulties scaling towards Big Data
- Problems include:
  - Slow
  - Scalability issues
  - Unnecessary overhead
  - Poor support for unstructured data
- New technologies emerge to resolve these problems ⇒ NoSQL

Non-Relational DB

- **Non-relational DB** are also known as "**NoSQL DBs**"
- NoSQL databases started gaining popularity in the 2000's when companies began investing and researching more into distributed databases
- NoSQL DBs represent data in collections of JSON documents E.g. MongoDB
- NoSQL DBs
  - No predefined schema
  - Records can have different fields as necessary (dynamic schema)

All in the NoSQL Family

NoSQL databases are geared toward managing large sets of varied and frequently updated data, often in distributed systems or the cloud. They avoid the rigid schemas associated with relational databases. But the architectures themselves vary and are separated into four primary classifications, although types are blending over time.

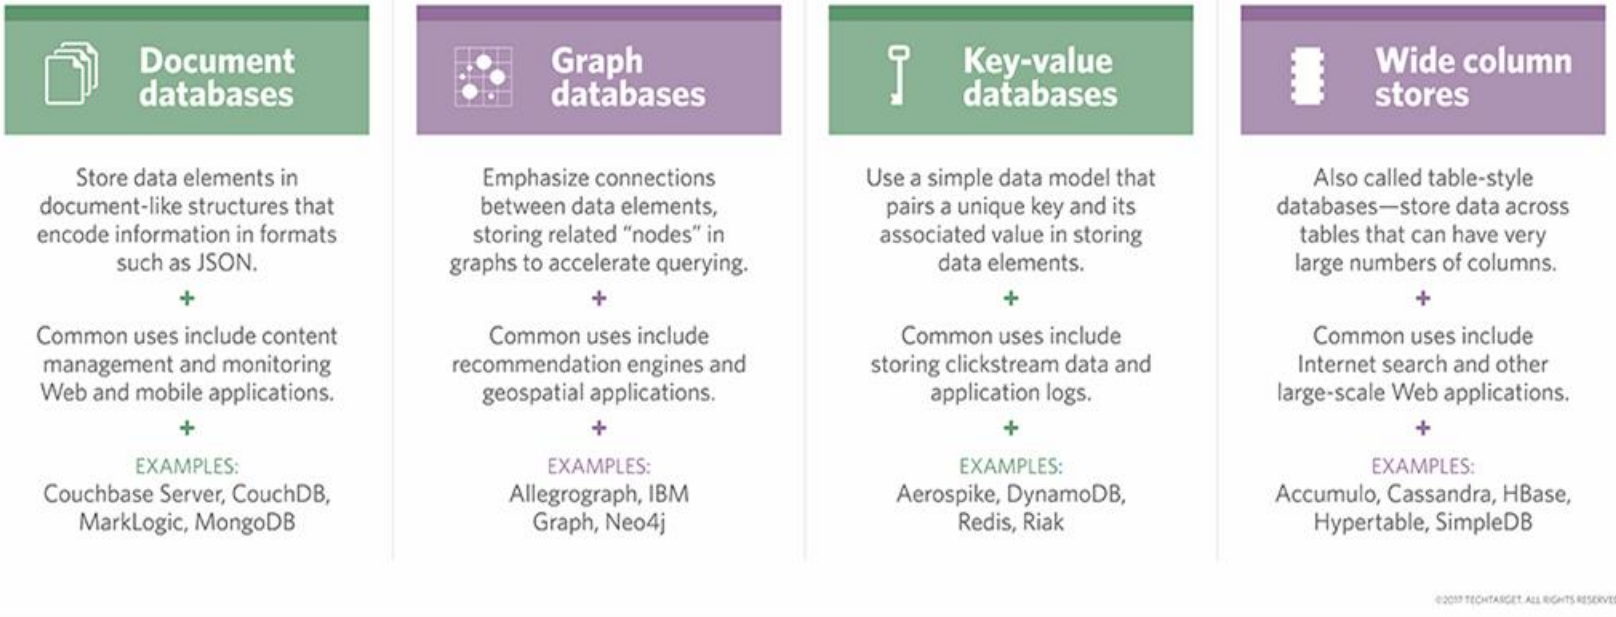

Source: https://searchdatamanagement.techtarget.com/definition/NoSQL-Not-Only-SQL

Key-value DBs

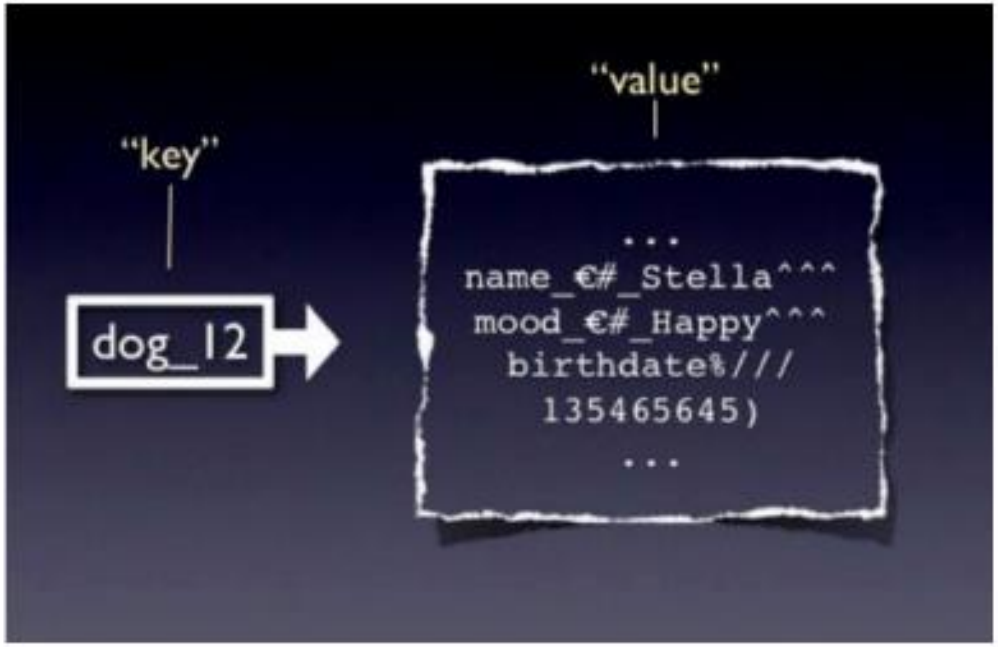

Key "dog_12": value_name "Stella", value_mood "Happy", etc

Column-based DBs

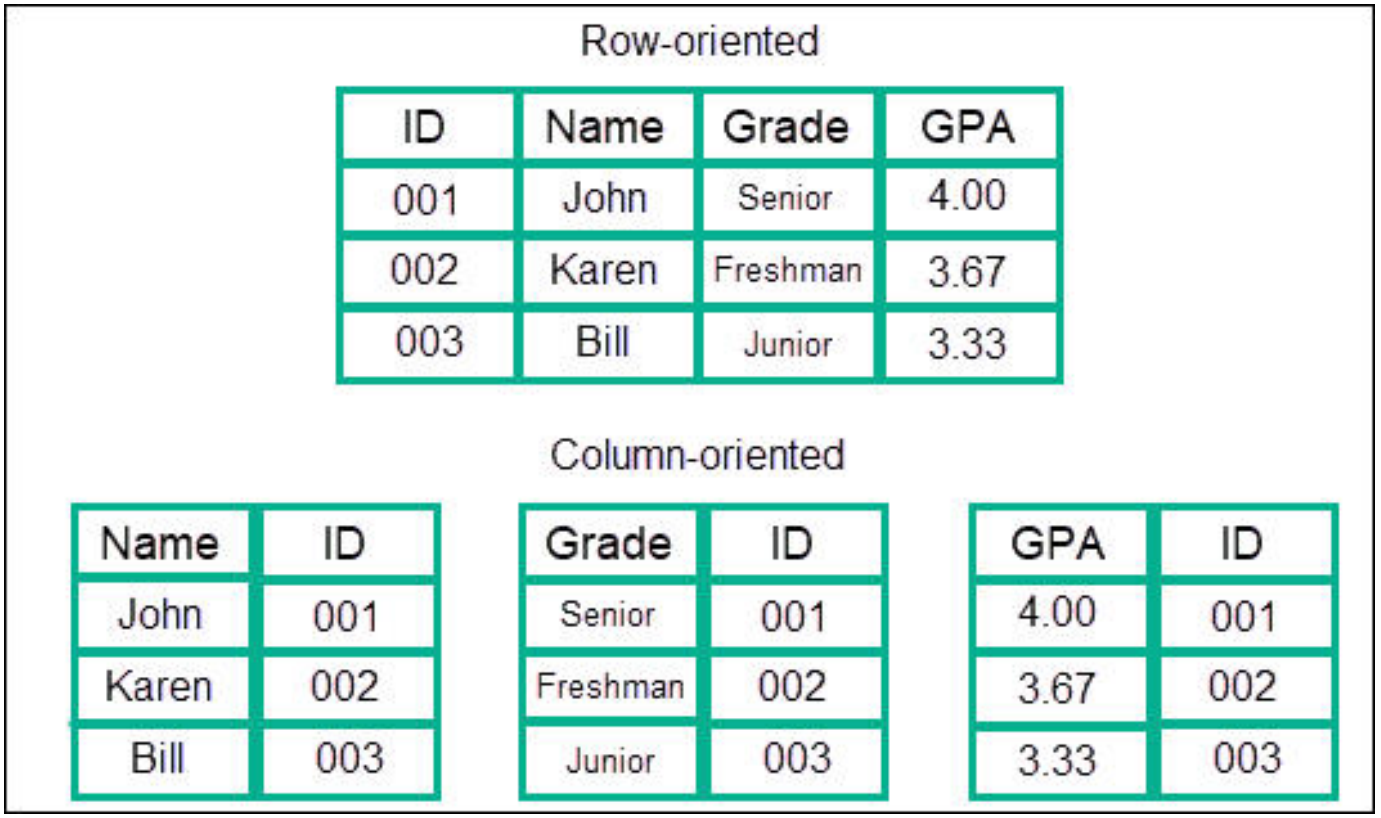

https://www.kdnuggets.com/2021/02/understanding-nosql-database-types-column-oriented-databases.html

Column-based DBs

Use the concept of keyspace. This keyspace contains all the column families, which then contain rows, which then contain columns.

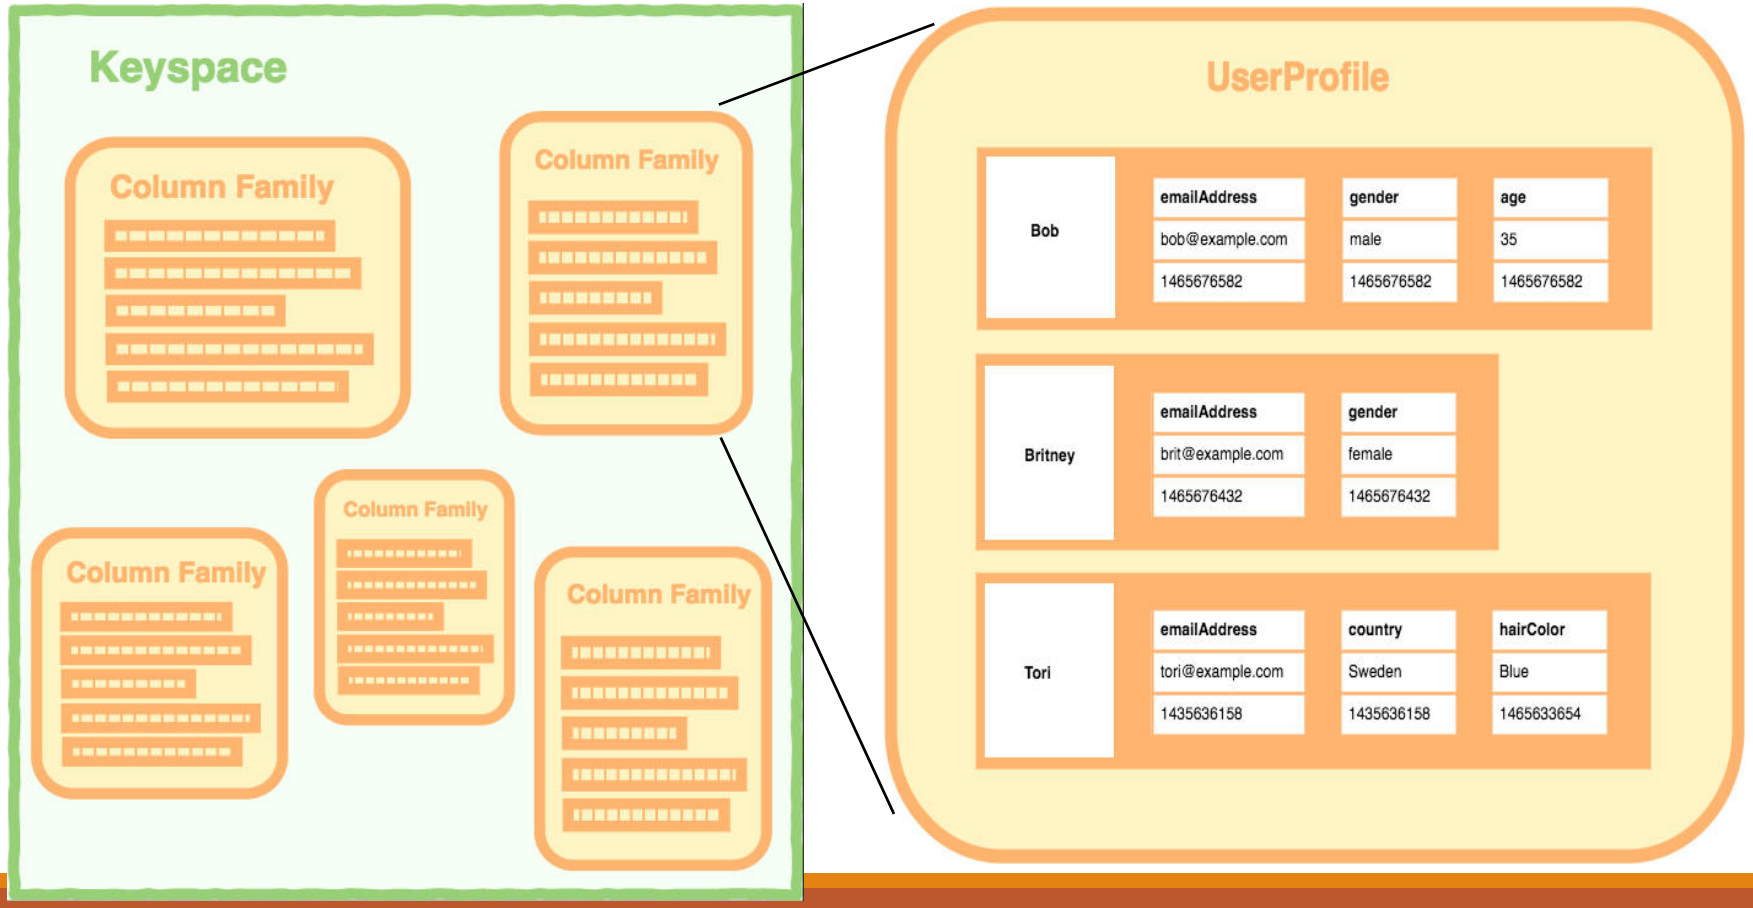

https://www.kdnuggets.com/2021/02/understanding-nosql-database-types-column-oriented-databases.html

Column-based DBs

Use the concept of keyspace. This keyspace contains all the column families, which then contain rows, which then contain columns.

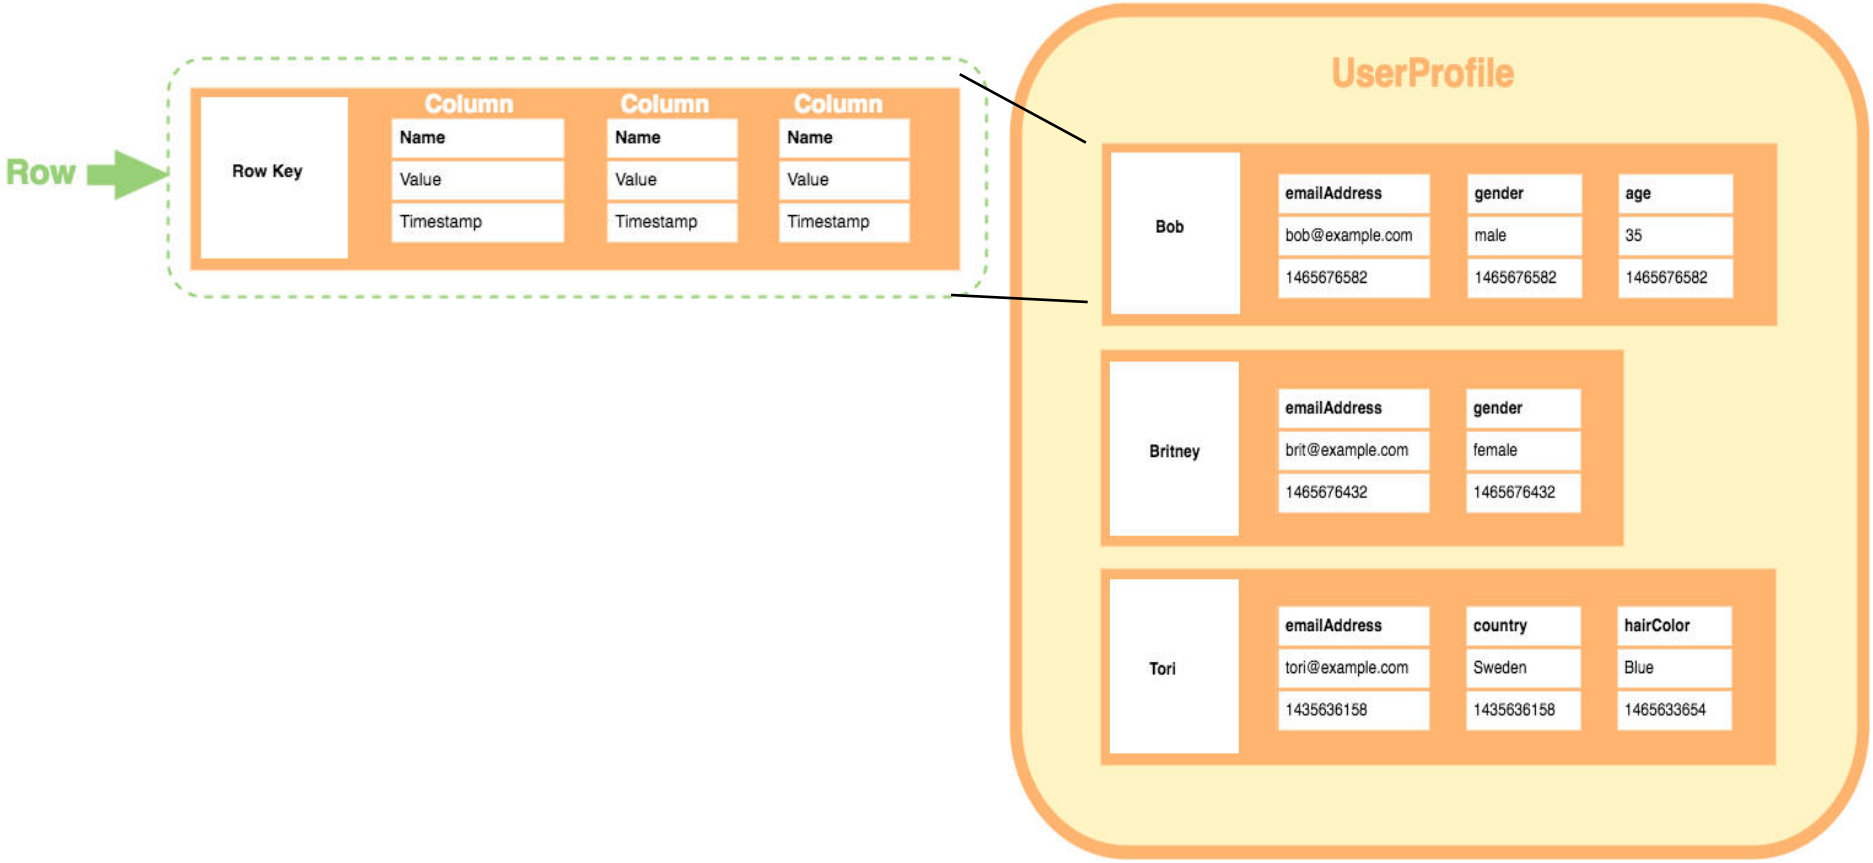

https://www.kdnuggets.com/2021/02/understanding-nosql-database-types-column-oriented-databases.html

Document-based DBs

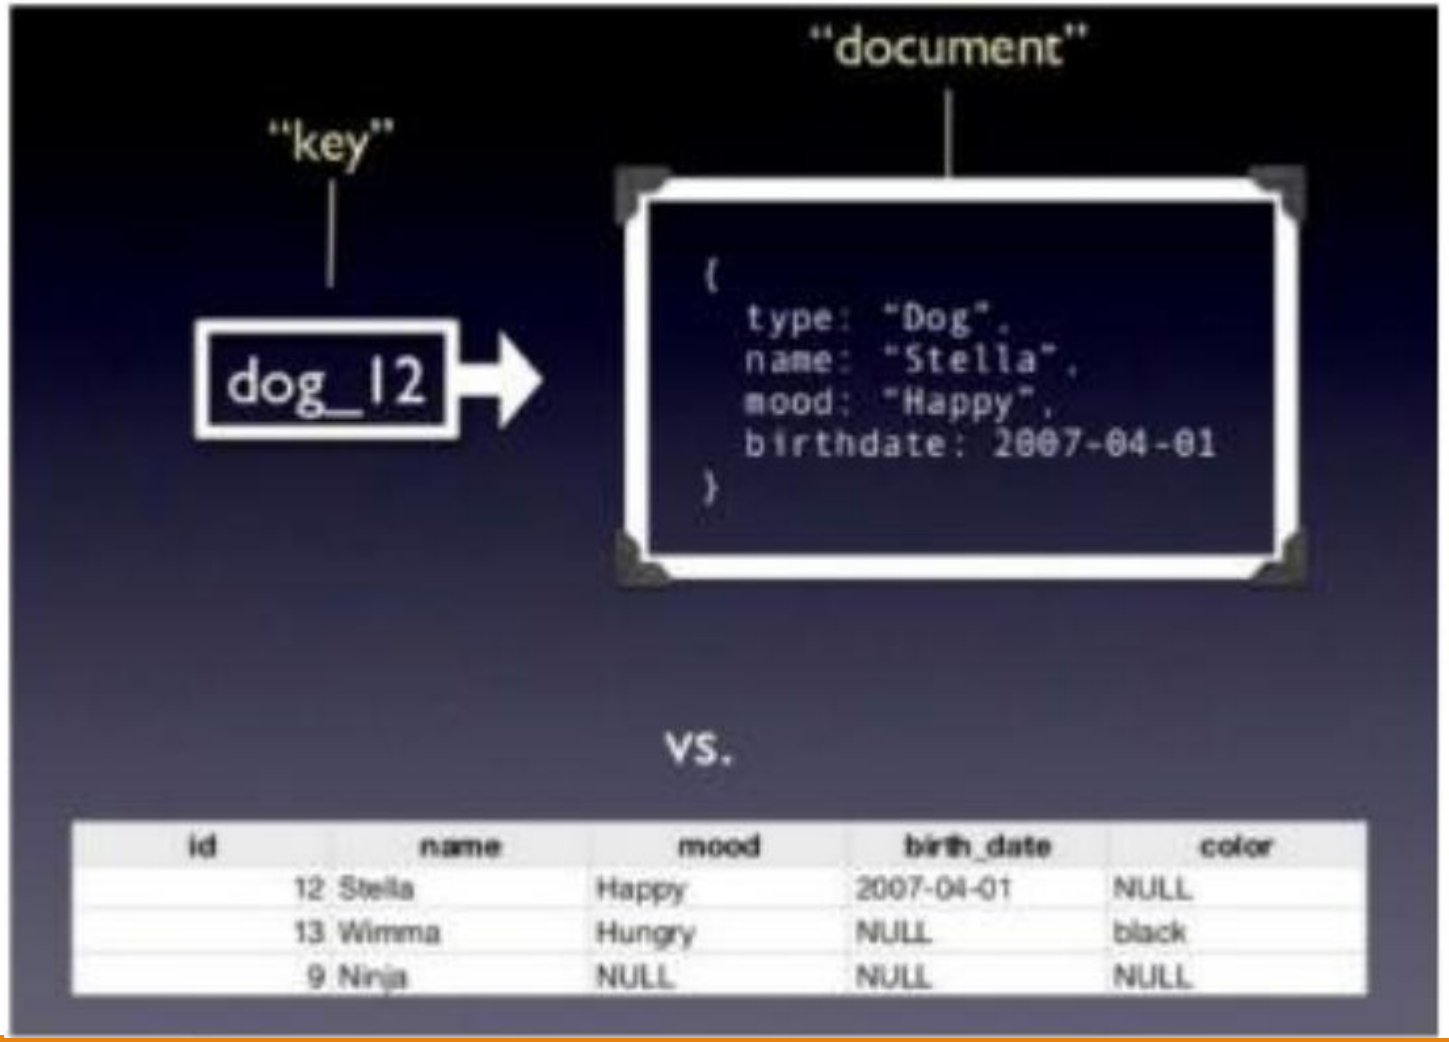

Document-based DBs

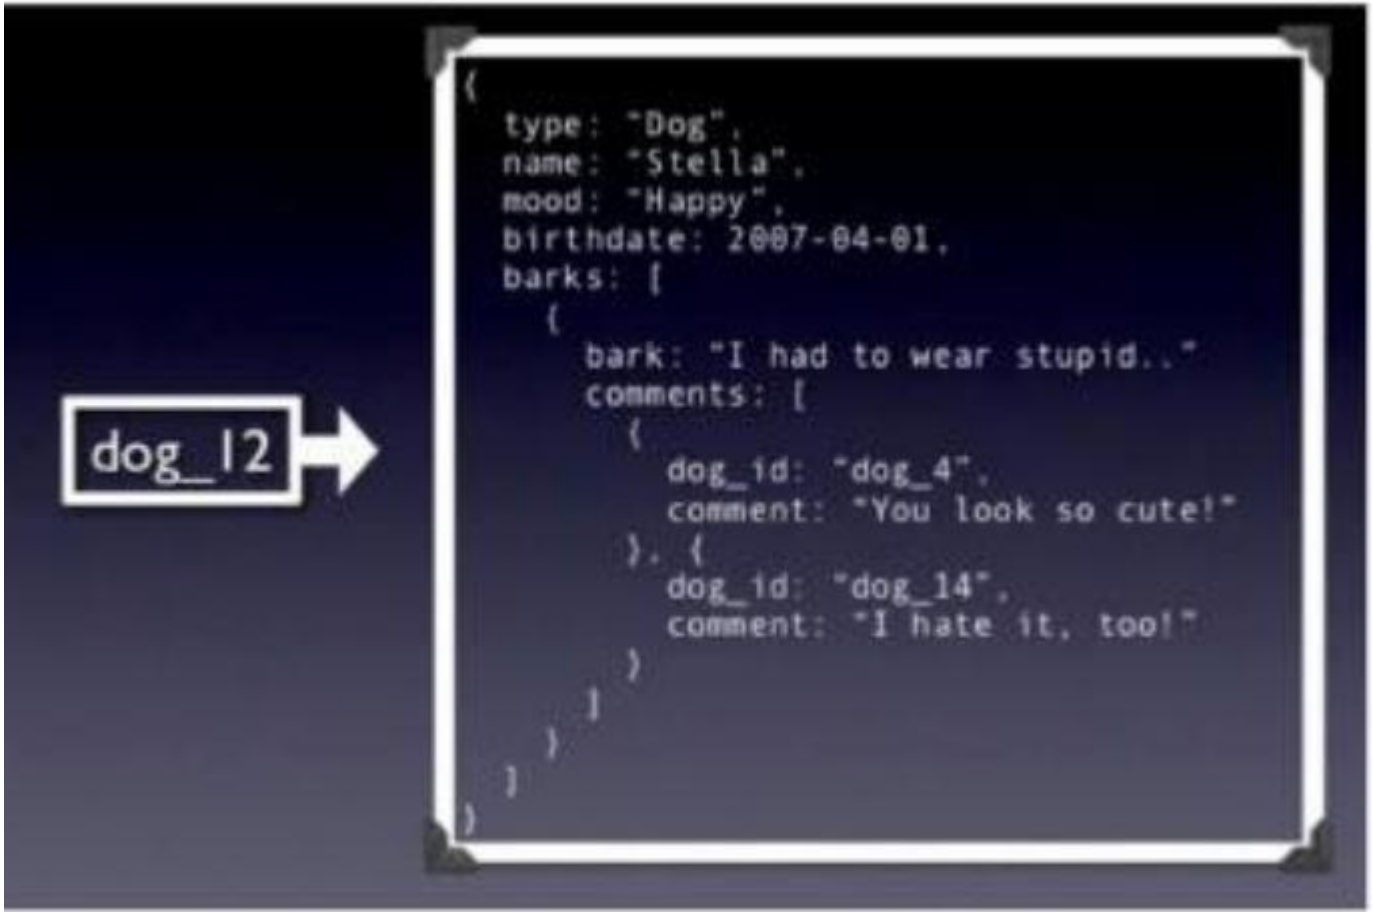

Graph DBs

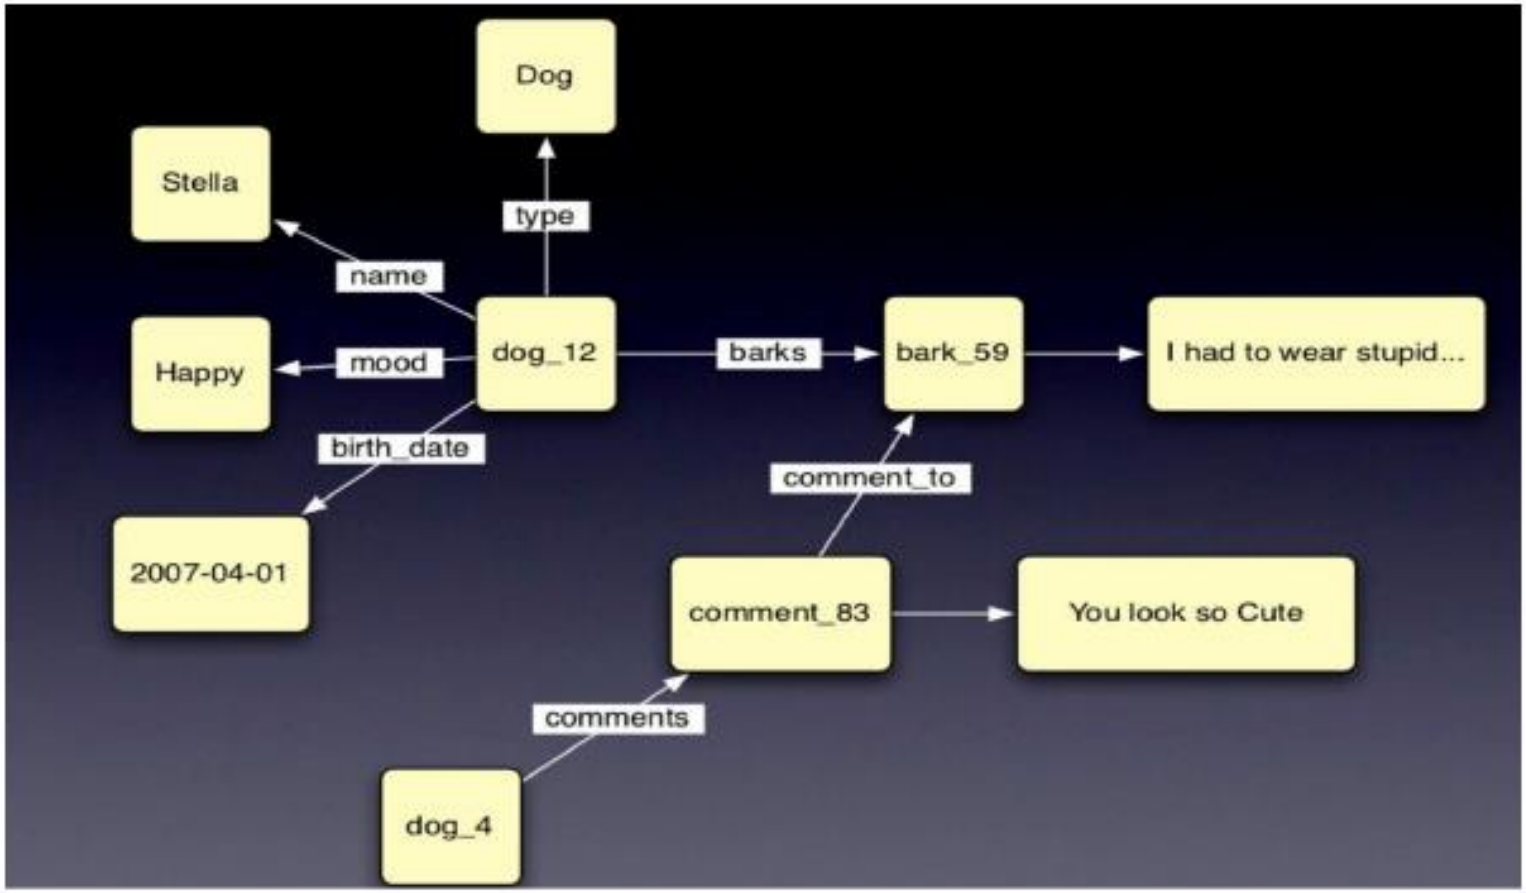
Dog
Stella
type
name
Happy
mood
dog_12
barks
bark_59
I had to wear stupid...
birth_date
2007-04-01
comment_to
comment_83
You look so Cute
comments
dog_4

Advantages of NoSQL DB

1. NoSQL databases generally process data faster than relational databases
2. NoSQL databases are also often faster because their data models are simpler
3. No schema required: Data can be inserted in a NoSQL database without first defining a rigid database schema. This provides immense flexibility, which ultimately delivers substantial business flexibility.

Comparison of DB Engines

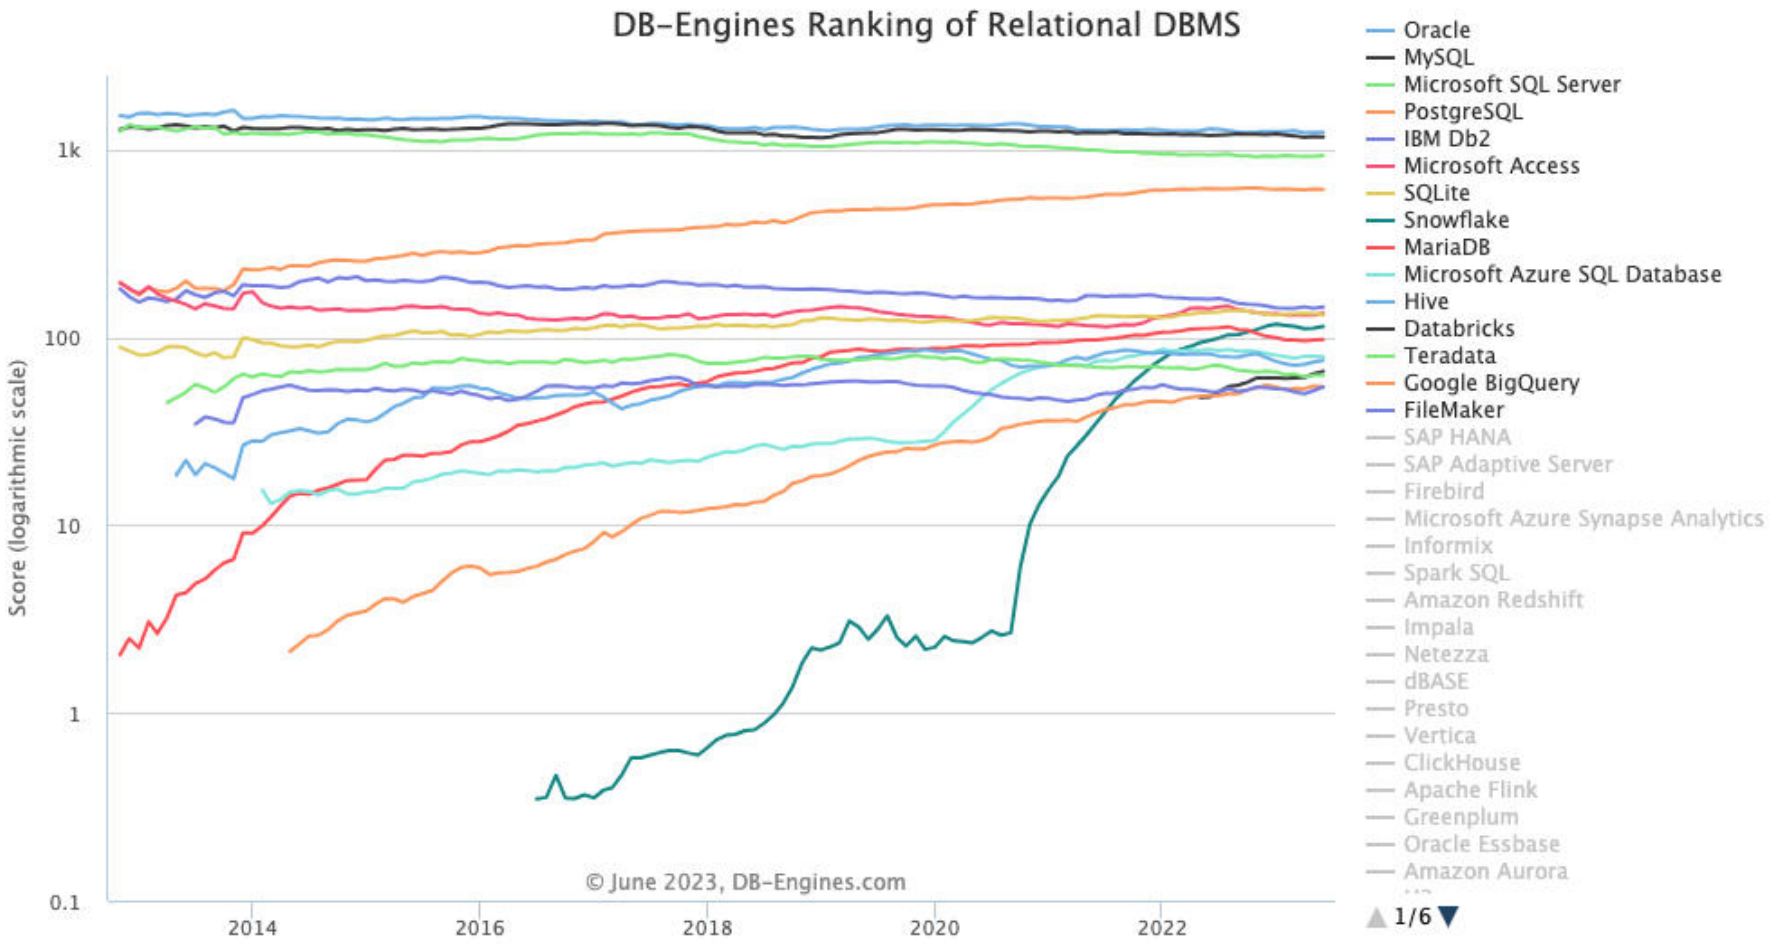

Source: https://www.linkedin.com/pulse/database-engine-ranking-popularity-2023-tamer-khraisha-ph-d-

Comparison of DB Engines

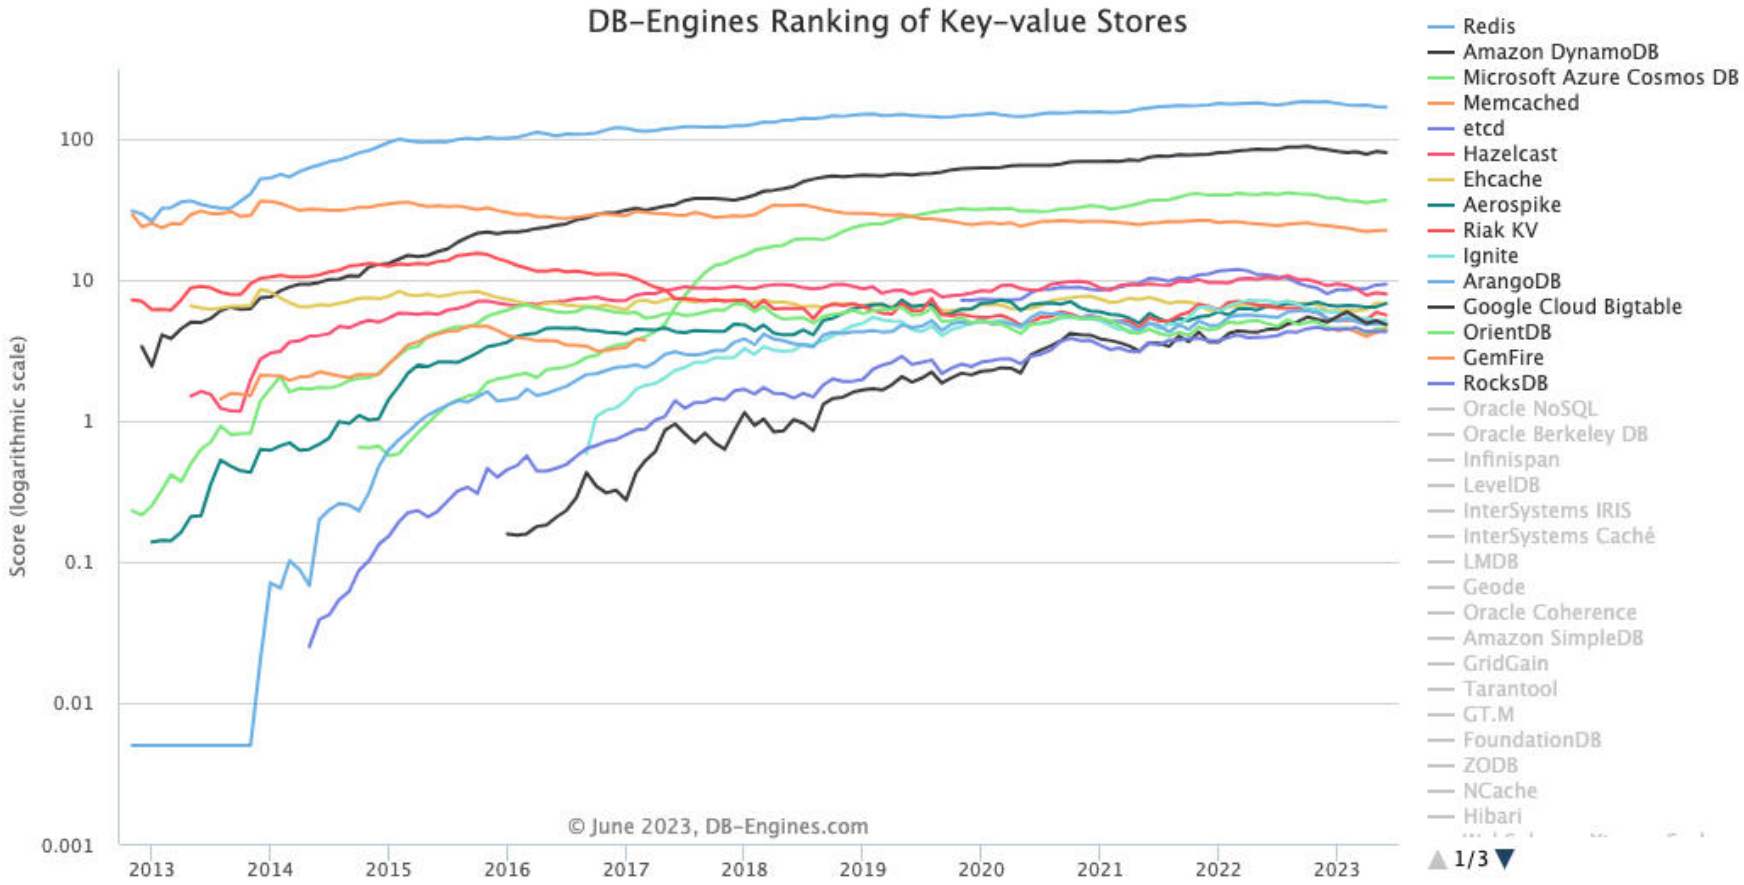

Source: https://www.linkedin.com/pulse/database-engine-ranking-popularity-2023-tamer-khraisha-ph-d-

Comparison of DB Engines

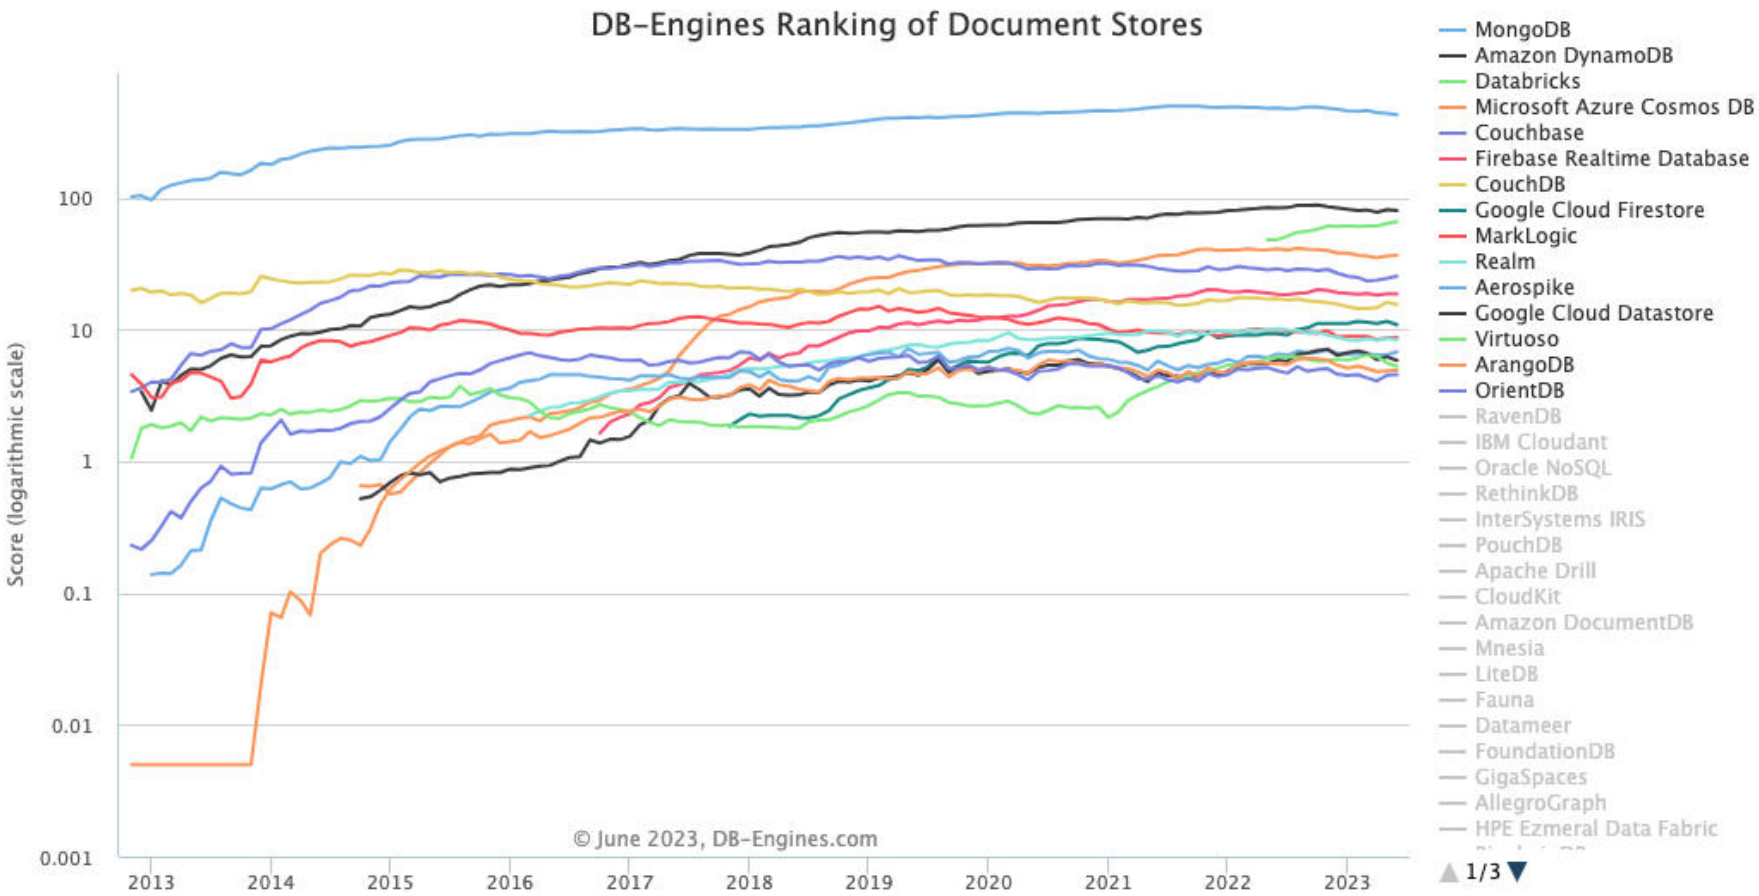

Source: https://www.linkedin.com/pulse/database-engine-ranking-popularity-2023-tamer-khraisha-ph-d-

Comparison of DB Engines

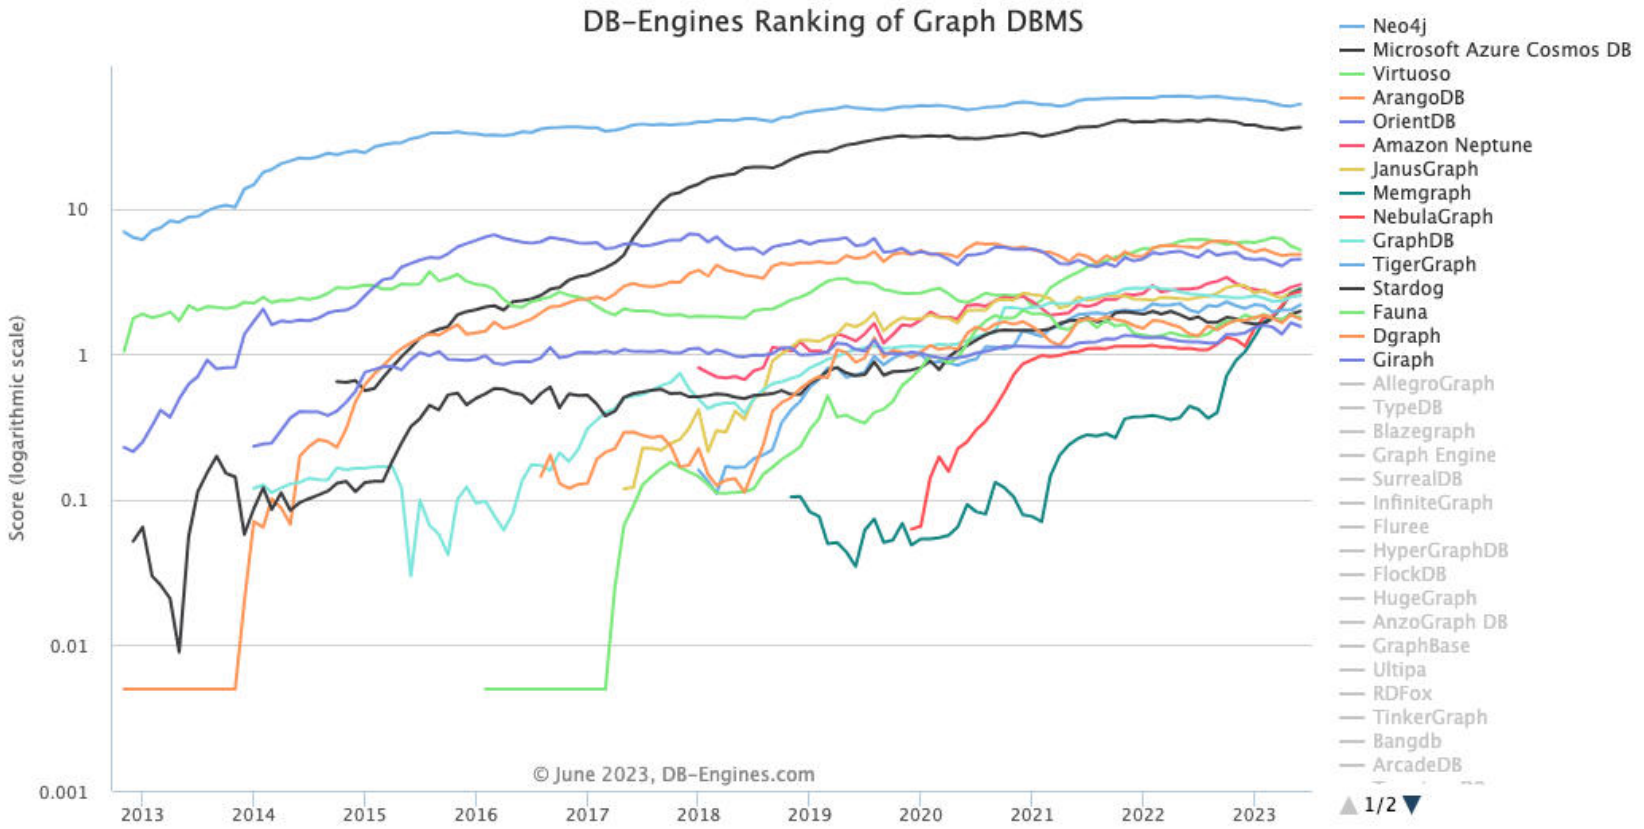

Source: https://www.linkedin.com/pulse/database-engine-ranking-popularity-2023-tamer-khraisha-ph-d-

End of Lecture 2